

Probabilistic Programming is great



Kristian Kersting



Alejandro Molina
TU Darmstadt



Robert Peharz
U. Cambridge



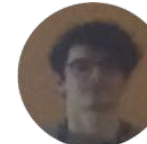
Zoubin Ghahramani
UBER AI Lab,
U. Cambridge



Martin Mladenov
Google Research



Sriraam Natarajan
UT Dallas



Antonio Vergari
MPI-IS



Isabell Valera
MPI-IS



Martin Grohe
RWTH Aachen



Amir Globerson
U. Haifa



Christopher Re
Stanford



Karl Stelzner
TU Darmstadt

... and many more

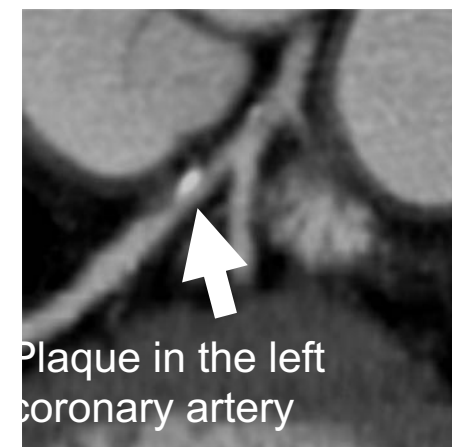
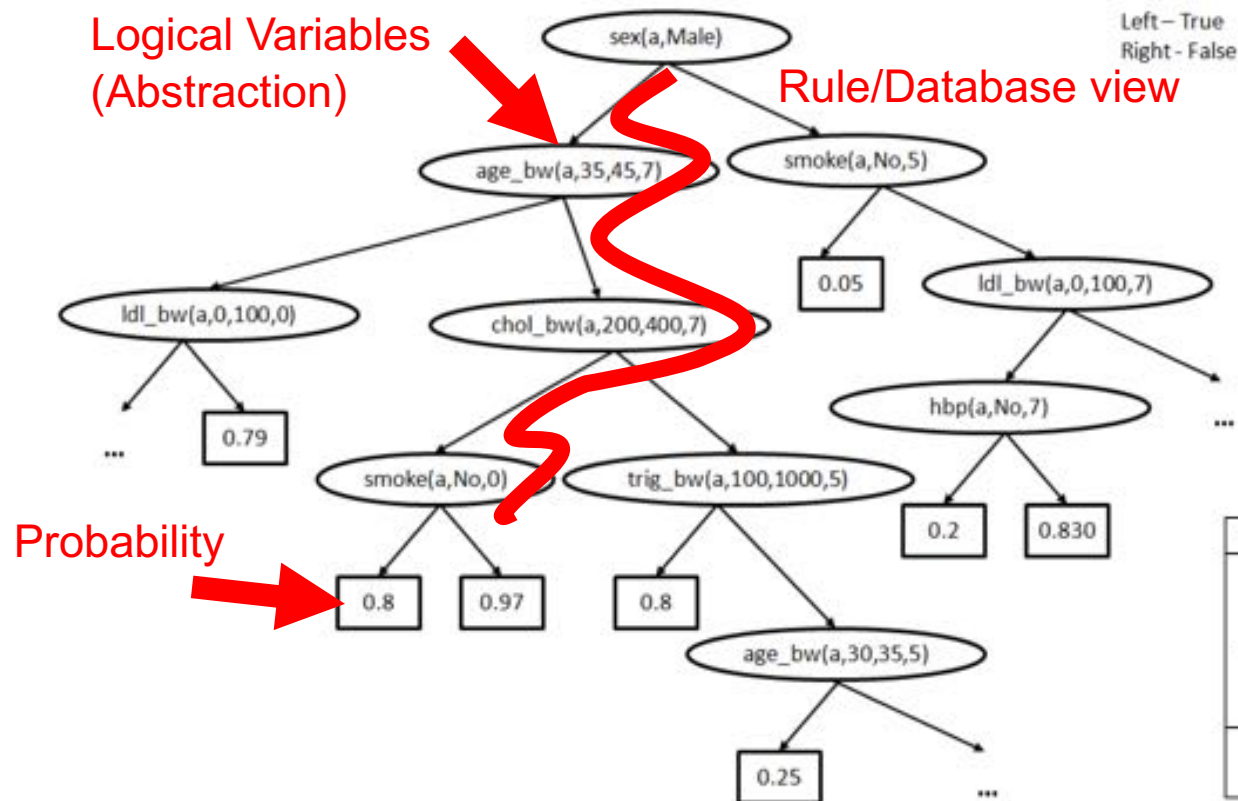
Probabilistic Programming is great

**But what if you have to condition on a
complete electronic health record (EHR)**



ProbProg over relational DBs

Atherosclerosis is the cause of the majority of Acute Myocardial Infarctions (heart attacks)



[Circulation; 92(8), 2157-62, 1995;
JACC; 43, 842-7, 2004]

Algorithm	Accuracy	AUC-ROC	The higher, the better
J48	0.667	0.607	
SVM	0.667	0.5	
AdaBoost	0.667	0.608	
Bagging	0.677	0.613	
NB	0.75	0.653	
RPT	0.669*	0.778] 25%
RFGB	0.667*	0.819	

Algorithm for Mining Markov Logic Networks	Likelihood The higher, the better	AUC-ROC The higher, the better	AUC-PR The higher, the better	Time The lower, the better
Boosting	0.81] 11%	0.96] 78%	0.93] 50%	9s] 37200x faster
LSM	0.73	0.54	0.62	93 hrs

[Kersting, Driessens ICML'08; Karwath, Kersting, Landwehr ICDM'08; Natarajan, Joshi, TadePELLi, Kersting, Shavlik. IJCAI'11; Natarajan, Kersting, Ip, Jacobs, Carr IAAI'13; Yang, Kersting, Terry, Carr, Natarajan AIME'15; Khot, Natarajan, Kersting, Shavlik ICDM'13, MLJ'12, MLJ'15]



Neural Information Processing Systems

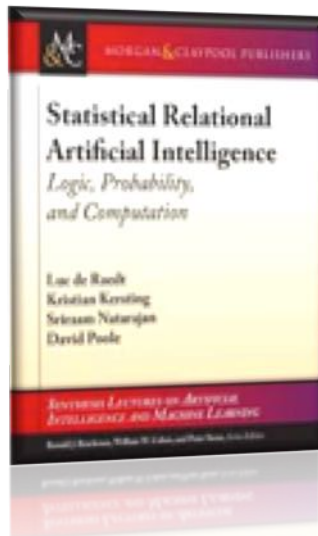
4. Dezember 2017 ·

NIPS 2017 Tutorial

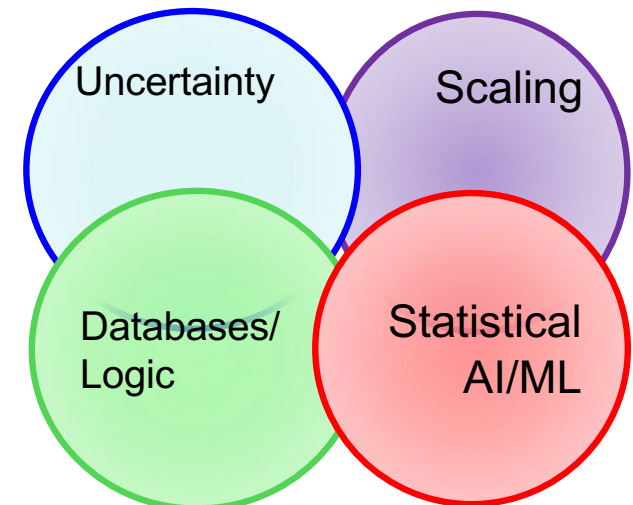
Now playing from NIPS 2017, Hall C, Statistical Relational Artificial Intelligence: Logic, Probability and Computation Tutorial.

3.075 Aufrufe

<https://www.facebook.com/nipsfoundation/videos/now-playing-from-nips-2017-hall-c-statistical-relational-artificial-intelligence/1552222671535633/>



De Raedt, Kersting, Natarajan, Poole: **Statistical Relational Artificial Intelligence: Logic, Probability, and Computation.** Morgan and Claypool Publishers, ISBN: 9781627058414, 2016.



<https://starling.utdallas.edu/software/boostsrl/wiki/>



People

Publications

Projects

Software

Datasets

Blog



BOOSTSRL BASICS

Getting Started

File Structure

Basic Parameters

Advanced Parameters

Basic Modes

Advanced Modes

ADVANCED BOOSTSRL

Default (RDN-Boost)

MLN-Boost

Regression

One-Class Classification

Cost-Sensitive SRL

Learning with Advice

Approximate Counting

Discretization of Continuous-Valued
Attributes

Lifted Relational Random Walks

Grounded Relational Random Walks

APPLICATIONS

Natural Language Processing

BoostSRL Wiki

BoostSRL (Boosting for Statistical Relational Learning) is a gradient-boosting based approach to learning different types of SRL models. As with the standard gradient-boosting approach, our approach turns the model learning problem to learning a sequence of regression models. The key difference to the standard approaches is that we learn relational regression models i.e., regression models that operate on relational data. We assume the data in a predicate logic format and the output are essentially first-order regression trees where the inner nodes contain conjunctions of logical predicates. For more details on the models and the algorithm, we refer to our book on this topic.

Sriraam Natarajan, Tushar Khot, Kristian Kersting and Jude Shavlik, Boosted Statistical Relational Learners: From Benchmarks to Data-Driven Medicine . SpringerBriefs in Computer Science, ISBN: 978-3-319-13643-1, 2015

Human-in-the-loop learning

Probabilistic Programming is great

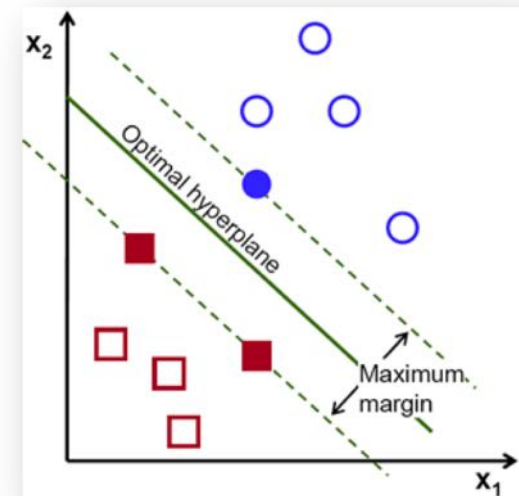
**But what if you want to feed a relational
DB into a Support Vector Machine?**

Standard SVM

$$\min_{\mathbf{w}, b, \xi} \mathcal{P}(\mathbf{w}, b, \xi) = \frac{1}{2} \mathbf{w}^2 + C \sum_{i=1}^n \xi_i$$

subject to $\begin{cases} \forall i & y_i(\mathbf{w}^\top \Phi(\mathbf{x}_i) + b) \geq 1 - \xi_i \\ \forall i & \xi_i \geq 0 \end{cases}$

Support Vector Machines
Cortes, Vapnik MLJ 20(3):273-297, 1995



High-level programming of QPs, using a few lines of code!

high-level AI / ML
programming languages

```
#QUADRATIC OBJECTIVE
minimize: sum{J in feature(I,J)} weight(J)**2 + c1 * slack + c2 * coslack;

#labeled examples should be on the correct side
subject to forall {I in labeled(I)}: labeled(I)*predict(I) >= 1 - slack(I);

#slacks are positive
subject to forall {I in labeled(I)}: slack(I) >= 0;

#TRANSDUCTIVE PART
#cited instances should have the same labels.
subject to forall {I1, I2 in linked(I1, I2)}: labeled(I1) * predict(I2) >= 1 - slack(I1, I2);
subject to forall {I1, I2 in linked(I1, I2)}: coslack(I1, I2) >= 0; #coslacks are positive
```

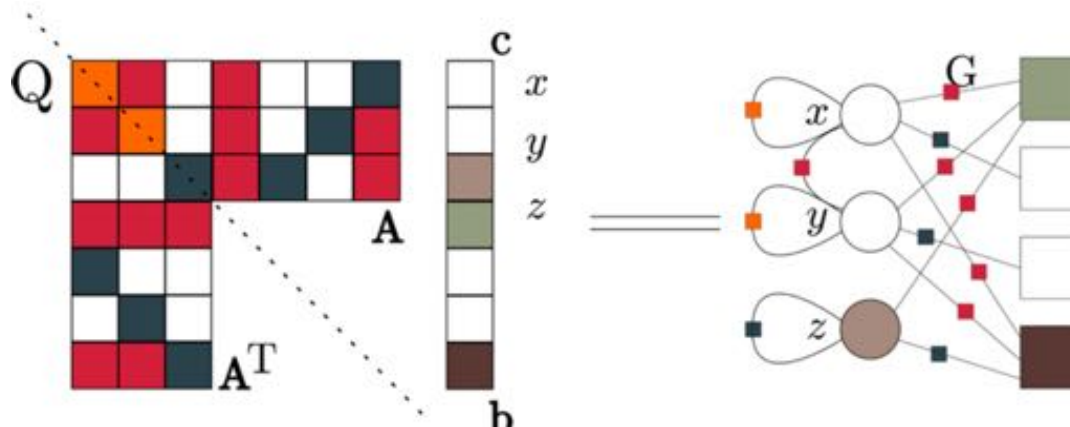
Citing papers should be on the same side of the hyperplane

Reduce the QP by running Weisfeiler-Lehman in quasi-linear time

$$\begin{aligned} \max_{[x,y,z]^T \in \mathbb{R}^3} \quad & 0x + 0y + 1z \\ \text{s.t.} \quad & -1z^2 - 2x^2 - 2y^2 + 1xy + 1yx \end{aligned}$$

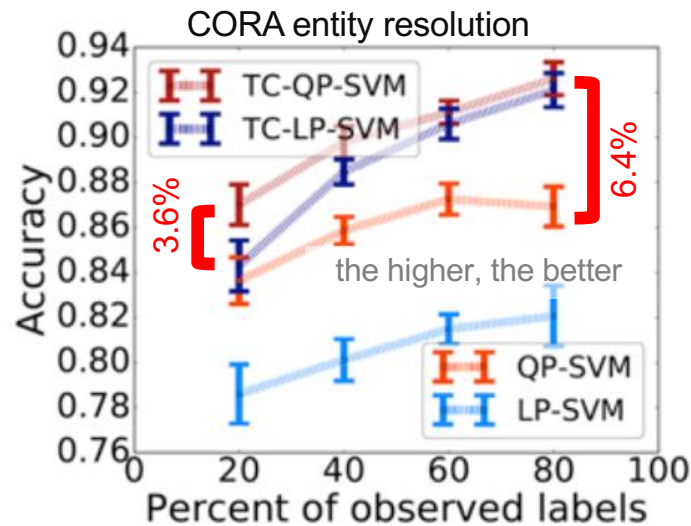
Compilers
for AI / ML
languages

$$\begin{bmatrix} 1 & 1 & 1 \\ -1 & 0 & 0 \\ 0 & -1 & 0 \\ 1 & 1 & -1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix} \leq \begin{bmatrix} 1 \\ 0 \\ 0 \\ -1 \end{bmatrix}$$



Statistical ML within relational DBs

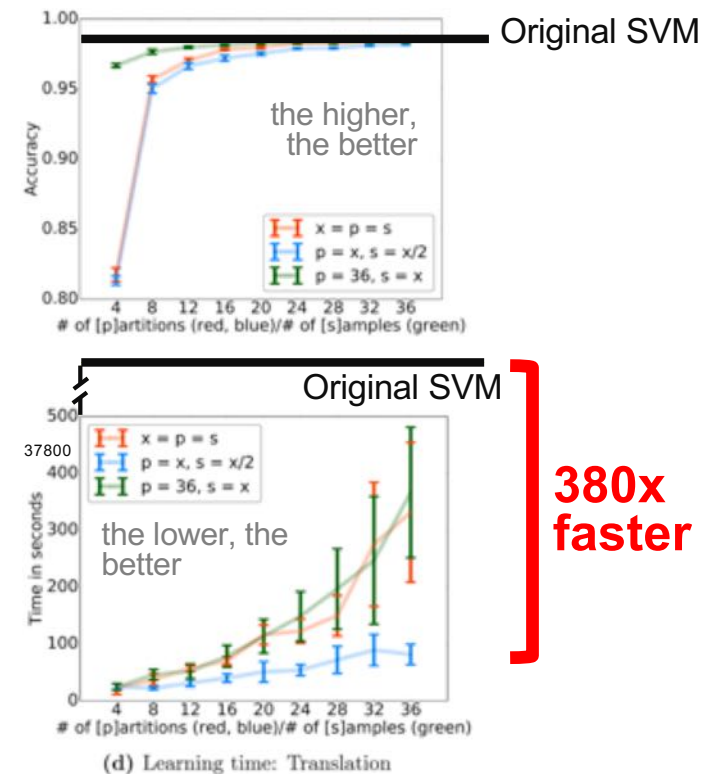
Collective Classification



On par with state-of-the-art by just few lines of code

Image Classification

MNIST image classification with label-preserving data programming



Faster than state-of-the-art



reloop

Overview

Source

Commits

Branches

Pull requests

Pipelines

Issues

Downloads

Reloop

1. Prerequisites as in requirements.txt

- Reloop requires Python 2.7+
- Scipy v0.15+
- Numpy v1.8.1+
- Cython v0.21.1+
- Cvxopt v1.1.7+
- Picos v1.0.1+
- infix v1.0.0+
- Ordered-Set v1.3.1+
- pyDatalog v0.14.6
- sympy v0.7.6+
- psycopg2 v2.6.1+
- problog v2.1.0.5+

Embedded within Python

reloop

RELOOP: A Toolkit for Relational Convex Optimization

<https://bitbucket.org/reloopdev/reloop>

If pip is available all prerequisites can be installed at once by running

```
$ pip install -r requirements.txt --upgrade
```

1.1 Optional Dependencies

These optional dependencies enable additional knowledge bases for usage. While Problog and SWI-Prolog both interface Prolog, psycopg2 interface a postgres database.

- Problog v2.1+
- Psycopg2 v2.6.1+
- SWI-Prolog

2. Installation

Once all the prerequisites have been installed simply run

```
python setup.py build_ext --inplace
```

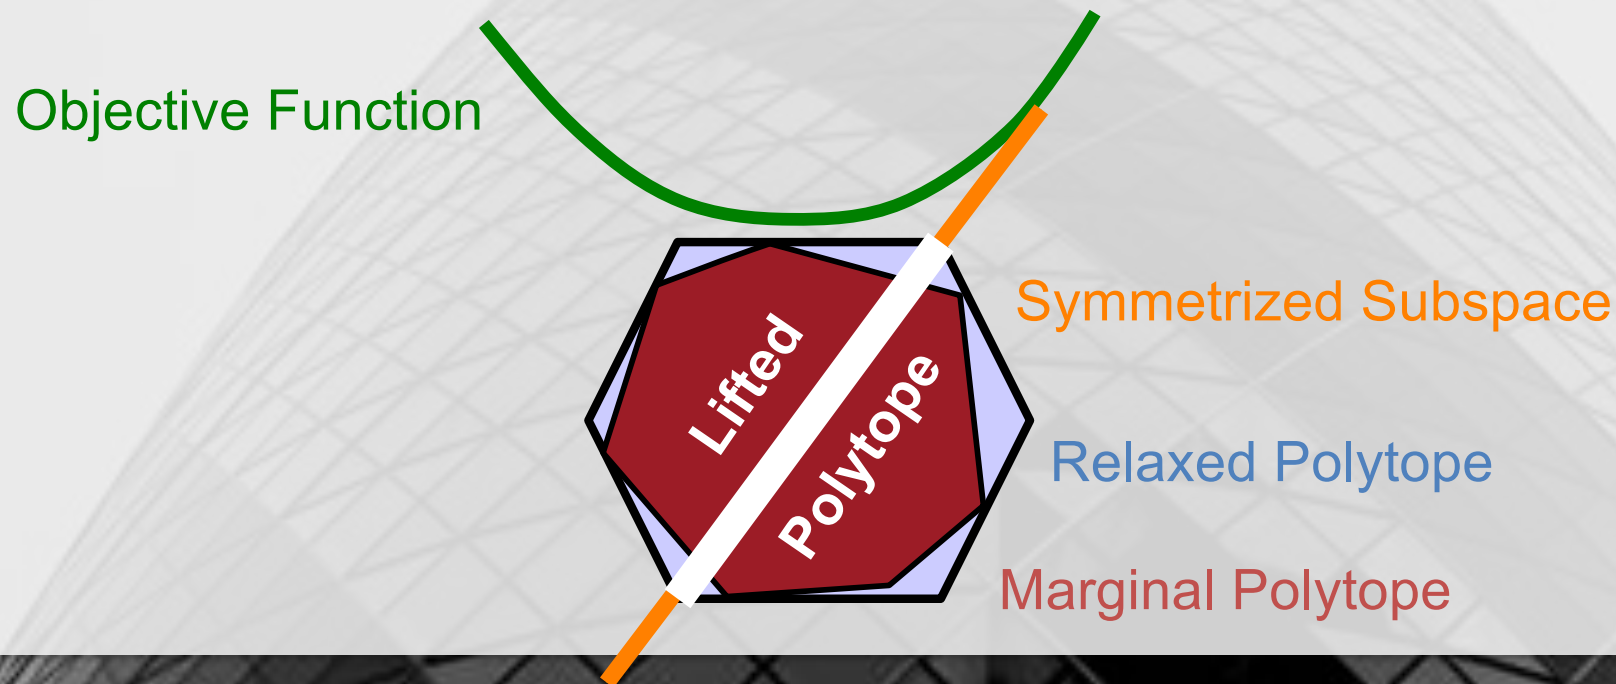
followed by either

```
python setup.py install
```

or

Algebraic approach for exploiting symmetries within probabilistic inference

$$\hat{\mathbf{x}} \in \arg \max_{\mathbf{x} \in \mathcal{X}^N} \left\{ \sum_{s \in V} \theta_s(x_s) + \sum_{(s,t) \in E} \theta_{st}(x_s, x_t) \right\}$$



**WL induces Fractional Automorphisms
of Linear and Quadratic Programs**

[Mladenov, Globerson, Kersting UAI '14, AISTATS '14, Mladenov, Kersting UAI '15]

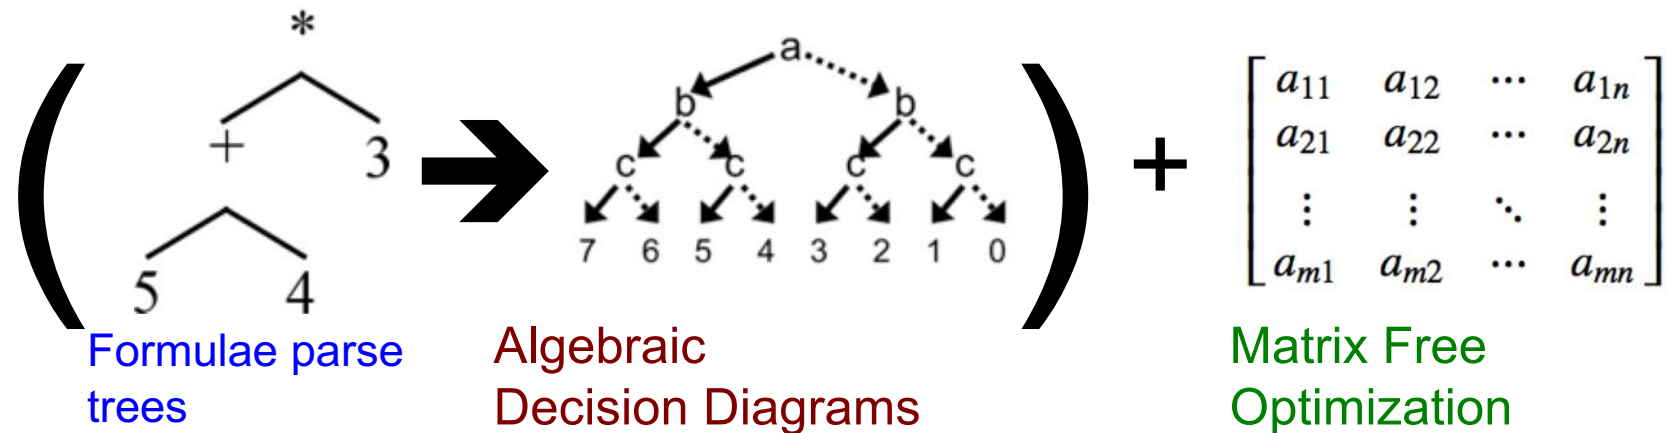
Algebraic approach for exploiting symmetries within probabilistic inference

Shows that all probabilistic inference approaches that relay on LPs/QPs such as MPLP, concave energies, ... can exploit symmetries

WL induces Fractional Automorphisms of Linear and Quadratic Programs

Other „-O“ flags lying ahead

New field: Symbolic-numerical AI



Problem Statistics				Symbolic IPM		Ground IPM
name	#vars	#constr	$nnz(A)$	ADD	time[s]	time[s]
factory	131.072	688.128	4.000.000	1819	6899	516
factory0	524.288	2.752.510	15.510.000	1895	6544	7920
factory1	2.097.150	11.000.000	59.549.700	2406	34749	159730
factory2	4.194.300	22.020.100	119.099.000	2504	36248	$\geq 48\text{hrs.}$

>4.8x faster

Applies to QPs but here illustrated on MDPs for a factory agent which must paint two objects and connect them. The objects must be smoothed, shaped and polished and possibly drilled before painting, each of which actions require a number of tools which are possibly available. Various painting and connection methods are represented, each having an effect on the quality of the job, and each requiring tools. Rewards (required quality) range from 0 to 10 and a discounting factor of 0.9 was used.

Probabilistic Programming is great

But what if you are not a statistician?

The Automatic Statistician

A system which explores an open-ended space of statistical models to discover a good explanation of the data, and then produces a detailed report with figures and natural-language text

...statistician probably, the team with a period of 10.8 years from 1643 and 1716 onwards.

...periodic with a period of 10.8 years. Across period ... smoothly with a typical lengthscale of 36.9 years. The shape of ... period is very smooth and resembles a sinusoid. This component applies ...

This component explains 71.5% of the residual variance; this increases the total variance from 77.8% to 92.3%. The addition of this component reduces the cross validated M from 0.18 to 0.15.

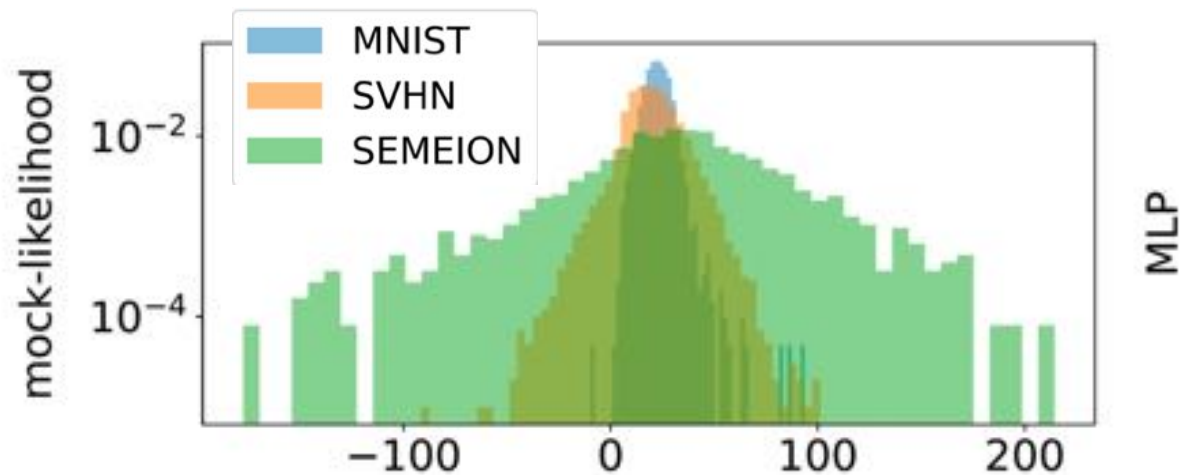
Only regression so far!



Llyod, Duvenaud, Ghahramani
U. Cambridge



Grosse, Tenenbaum
MIT



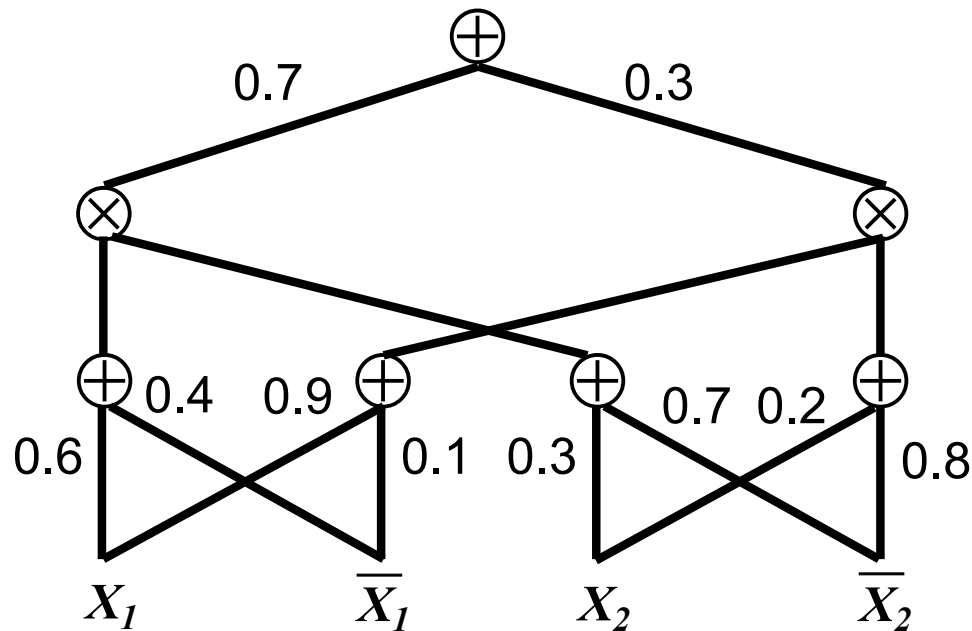
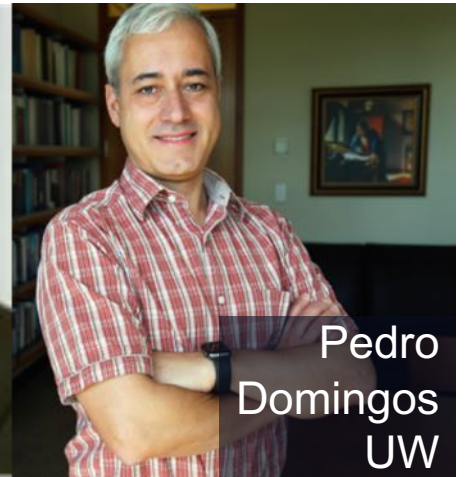
Deep neural networks may not be faithful probabilistic models

Deep Probabilistic Modelling using Sum-Product Networks

Adnan
Darwiche
UCLA



Pedro
Domingos
UW



Computational graph
(kind of TensorFlow
graphs) that encodes
how to compute
probabilities

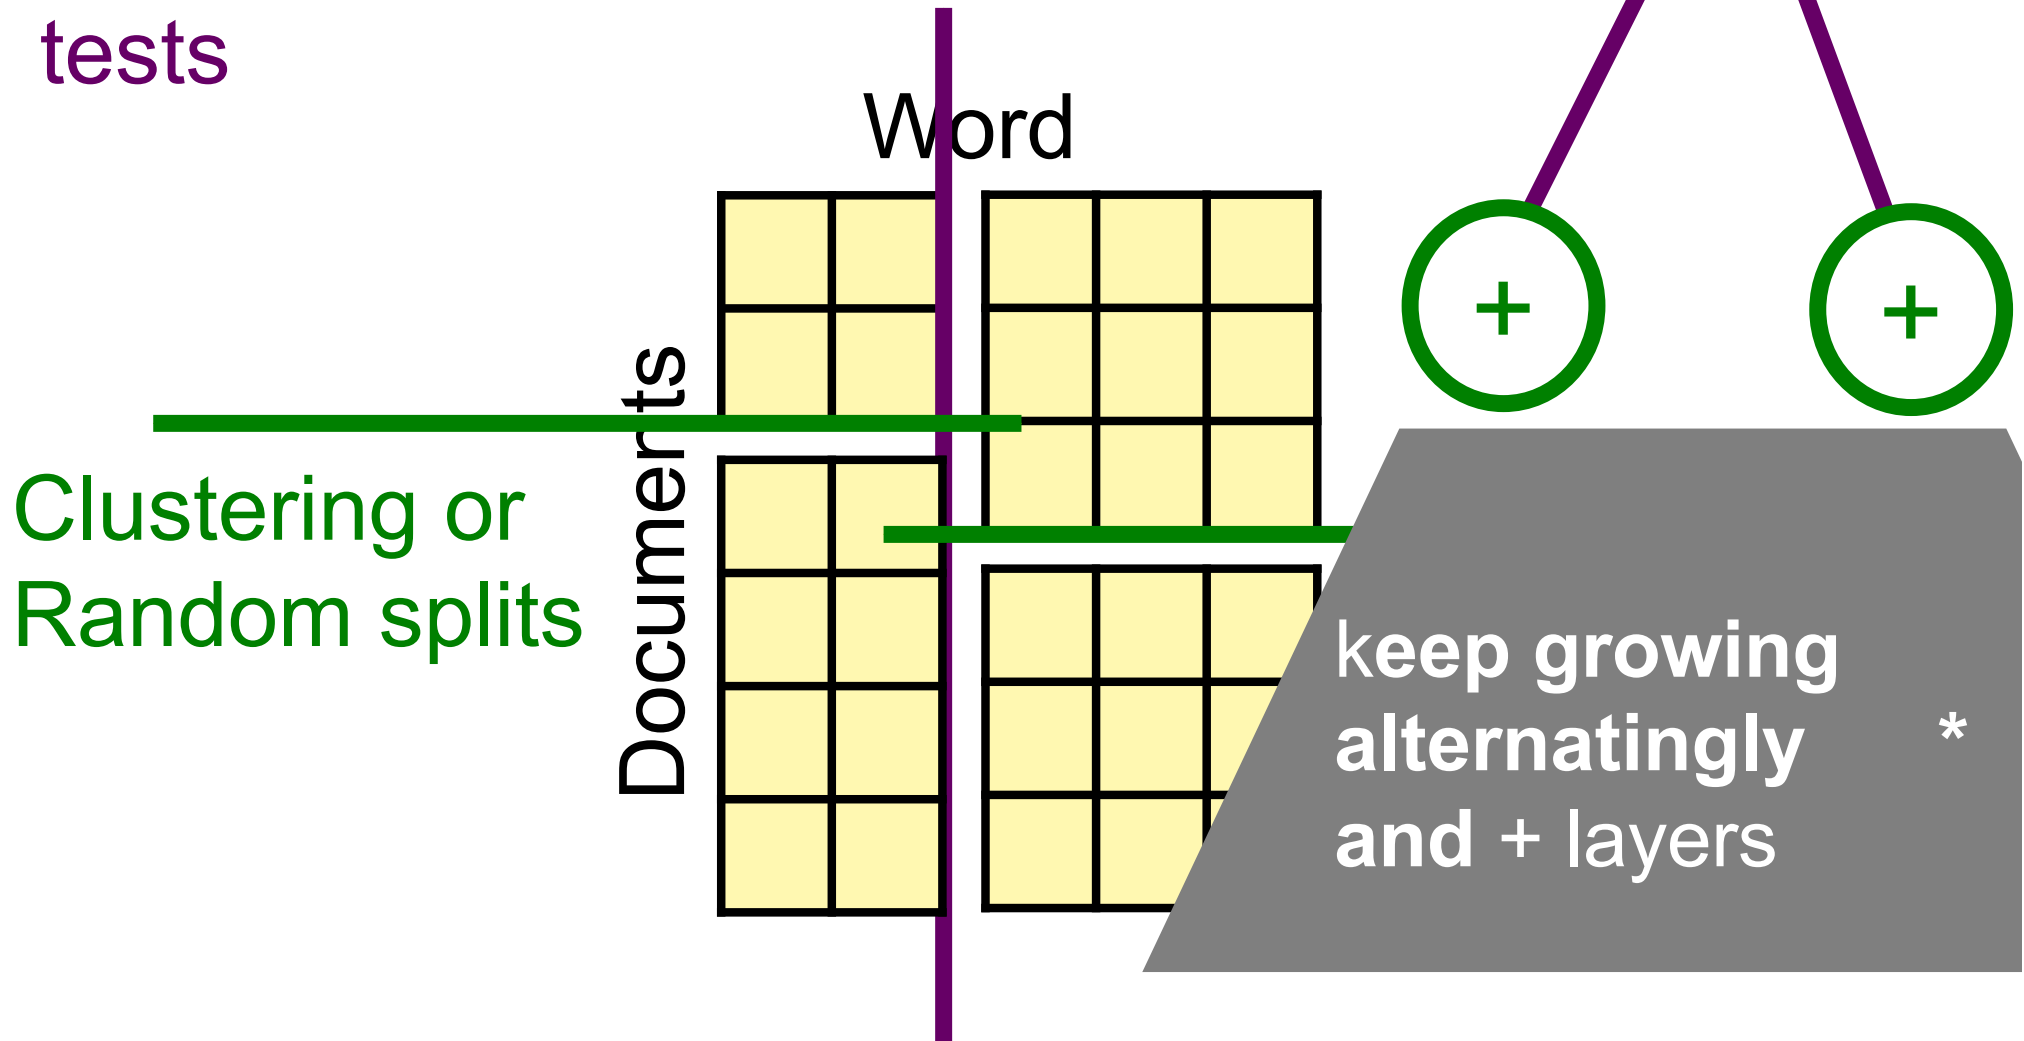
Inference is Linear in Size of Network

*SPNs are an instance of Arithmetic Circuits (ACs). ACs have been introduced into the AI literature more than 15 years ago as a tractable representation of probability distributions

[Darwiche CACM 48(4):608-647 2001]

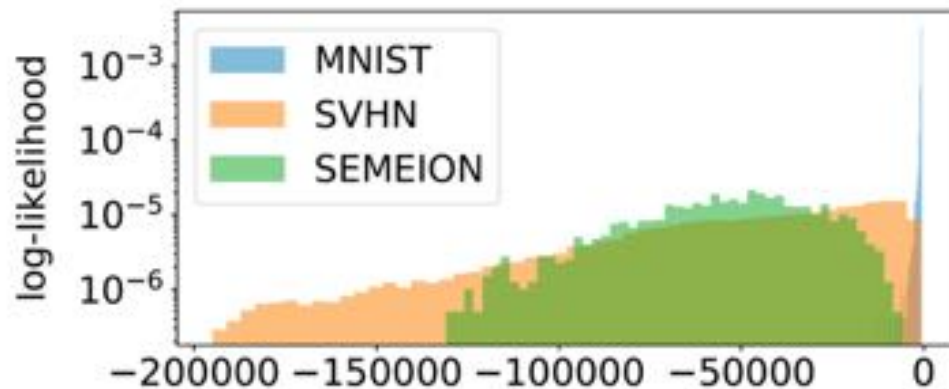
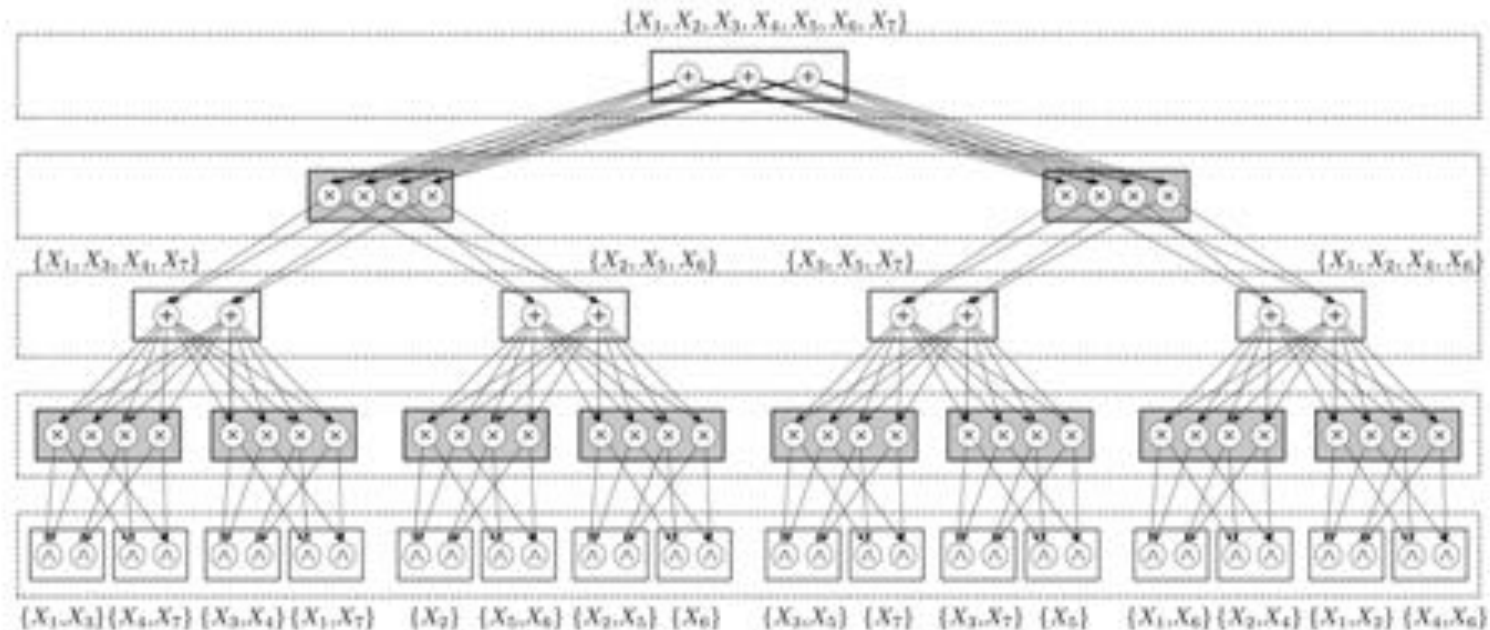
Greedy structure learning

Testing independence of random variables using e.g. nonparametric tests

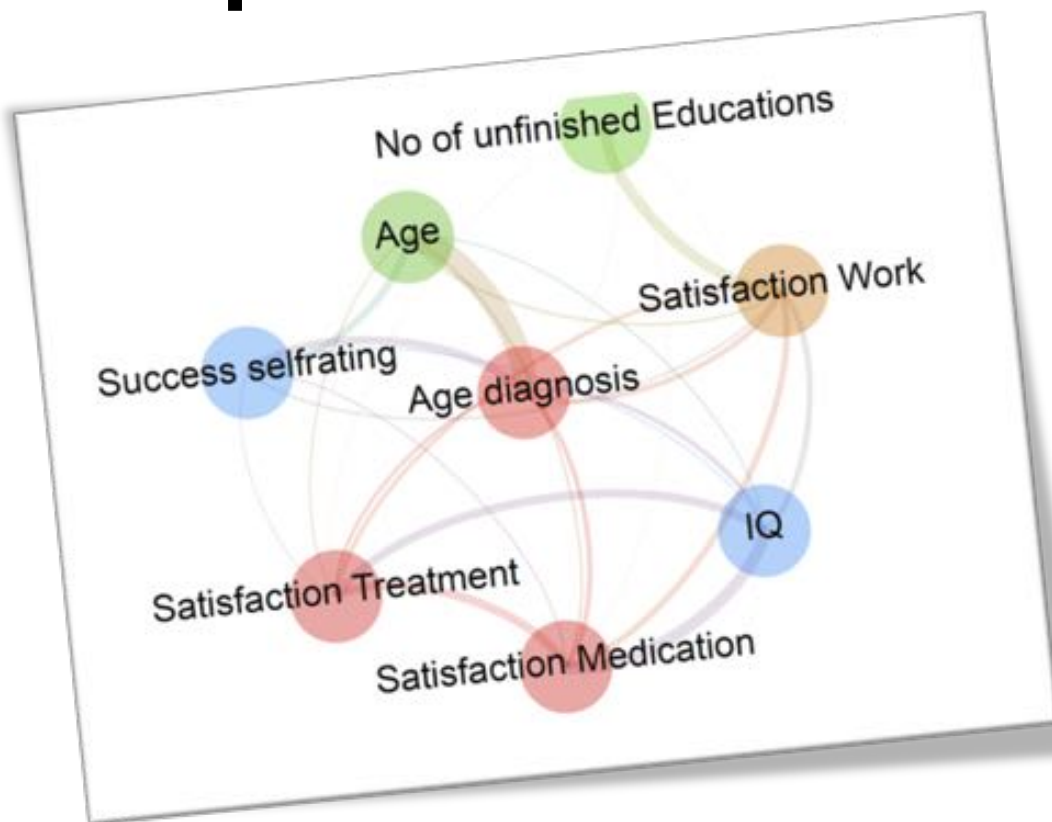


Random sum-product networks

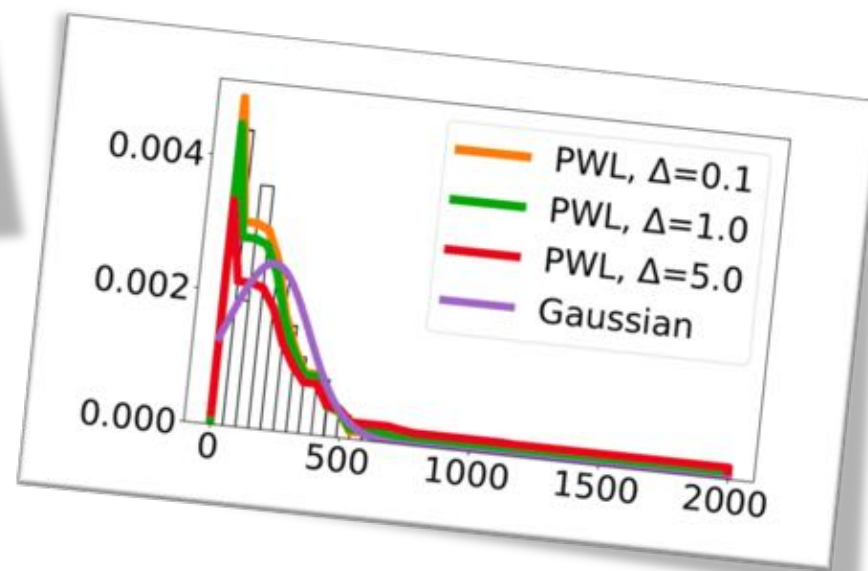
[Peharz, Vergari, Molina, Stelzner, Trapp, Kersting, Ghahramani UDL@UAI 2018]



Distribution-agnostic Deep Probabilistic Learning

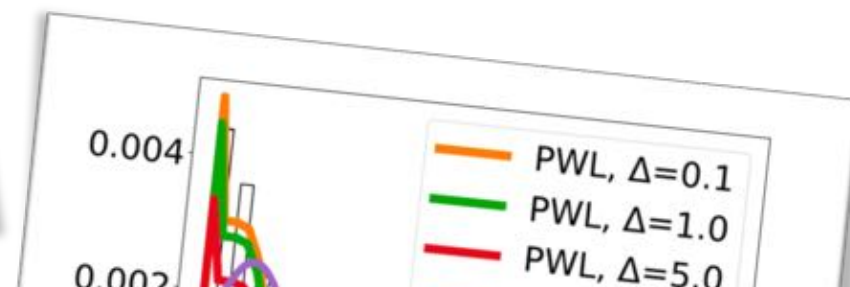
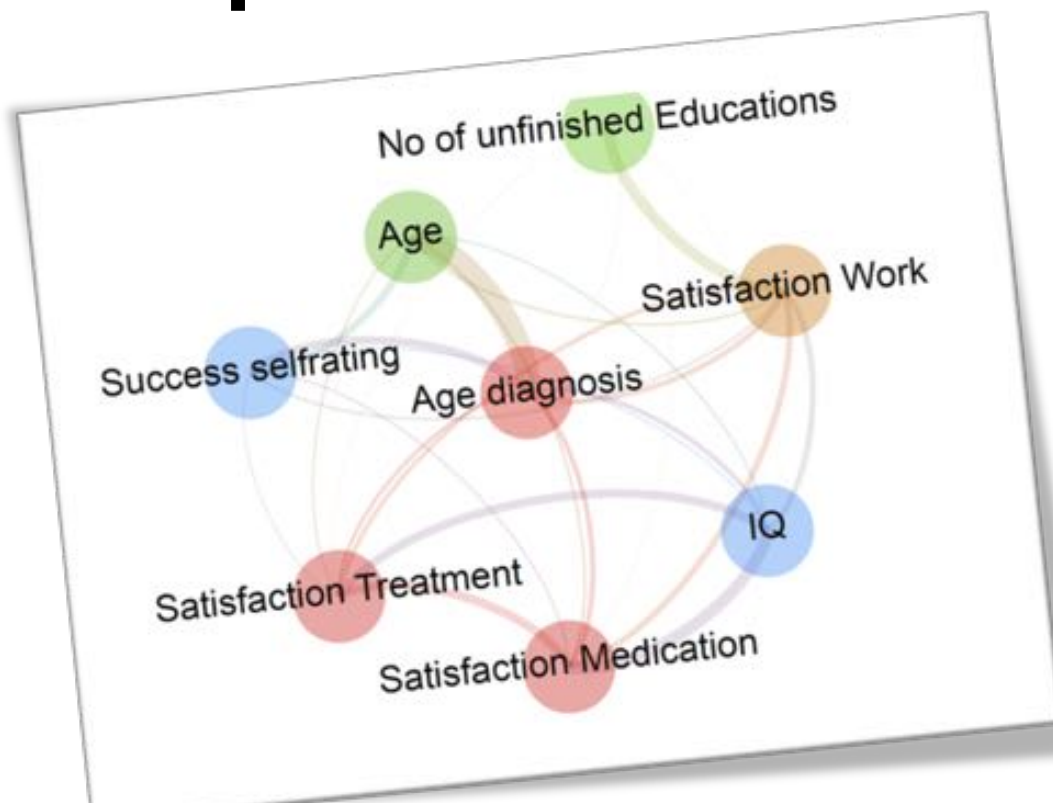


**Use nonparametric
independency tests
and piece-wise linear
approximations**



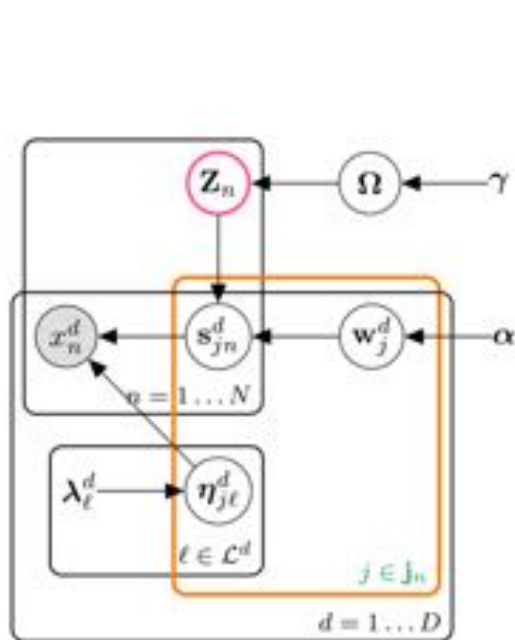
Distribution-agnostic Deep Probabilistic Learning

**Use nonparametric
independency tests
and piece-wise linear
approximations**

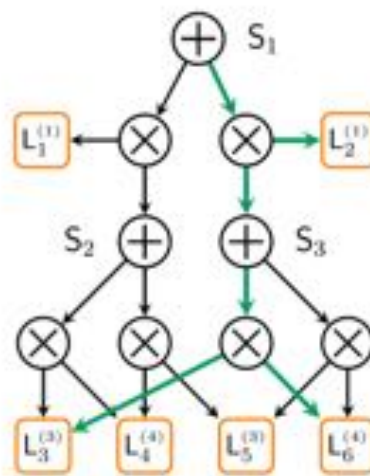


However, we have to provide the statistical types and do not gain insights into the parametric forms of the variables.
Are they Gaussians? Gammas? ...

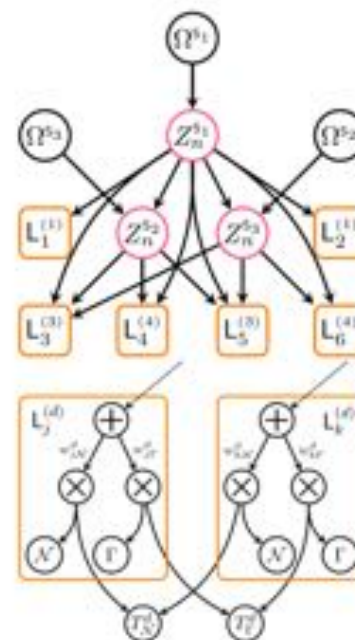
Automatic Bayesian Density Analysis



(a) Graphical model



(b) SPN



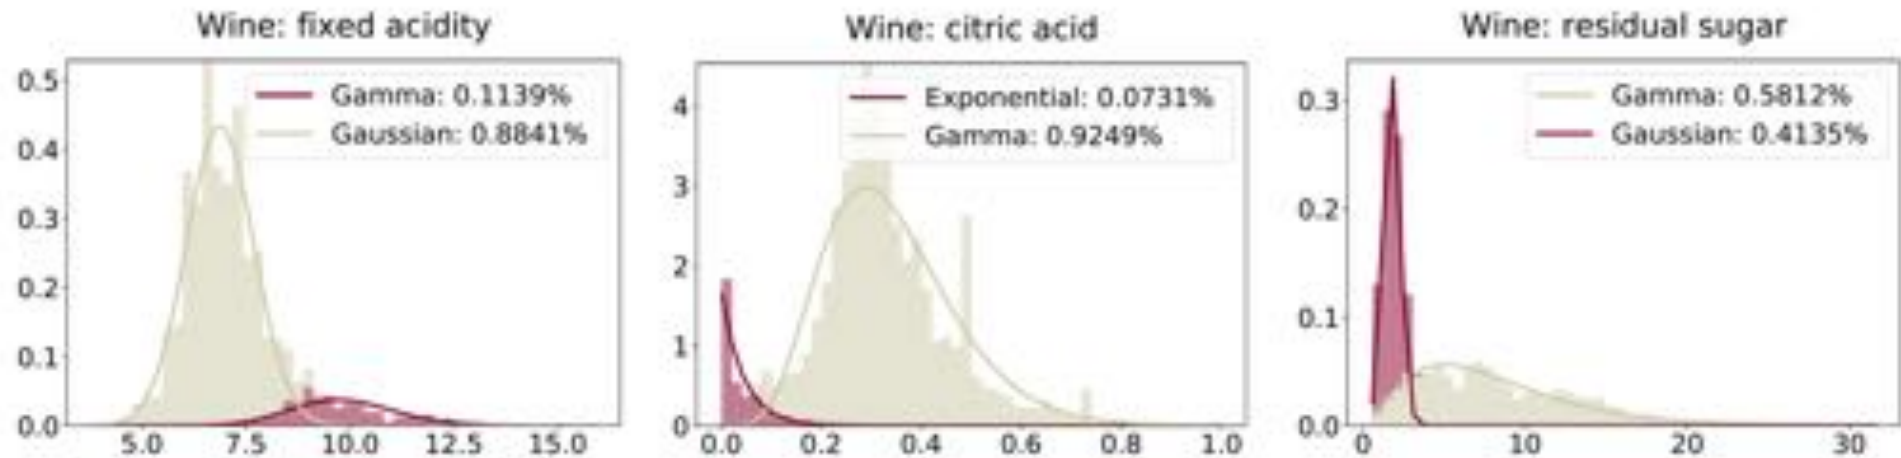
(c) Type-augmented SPN

**Bayesian discovery of
statistical types and
parametric forms of
variables**

+

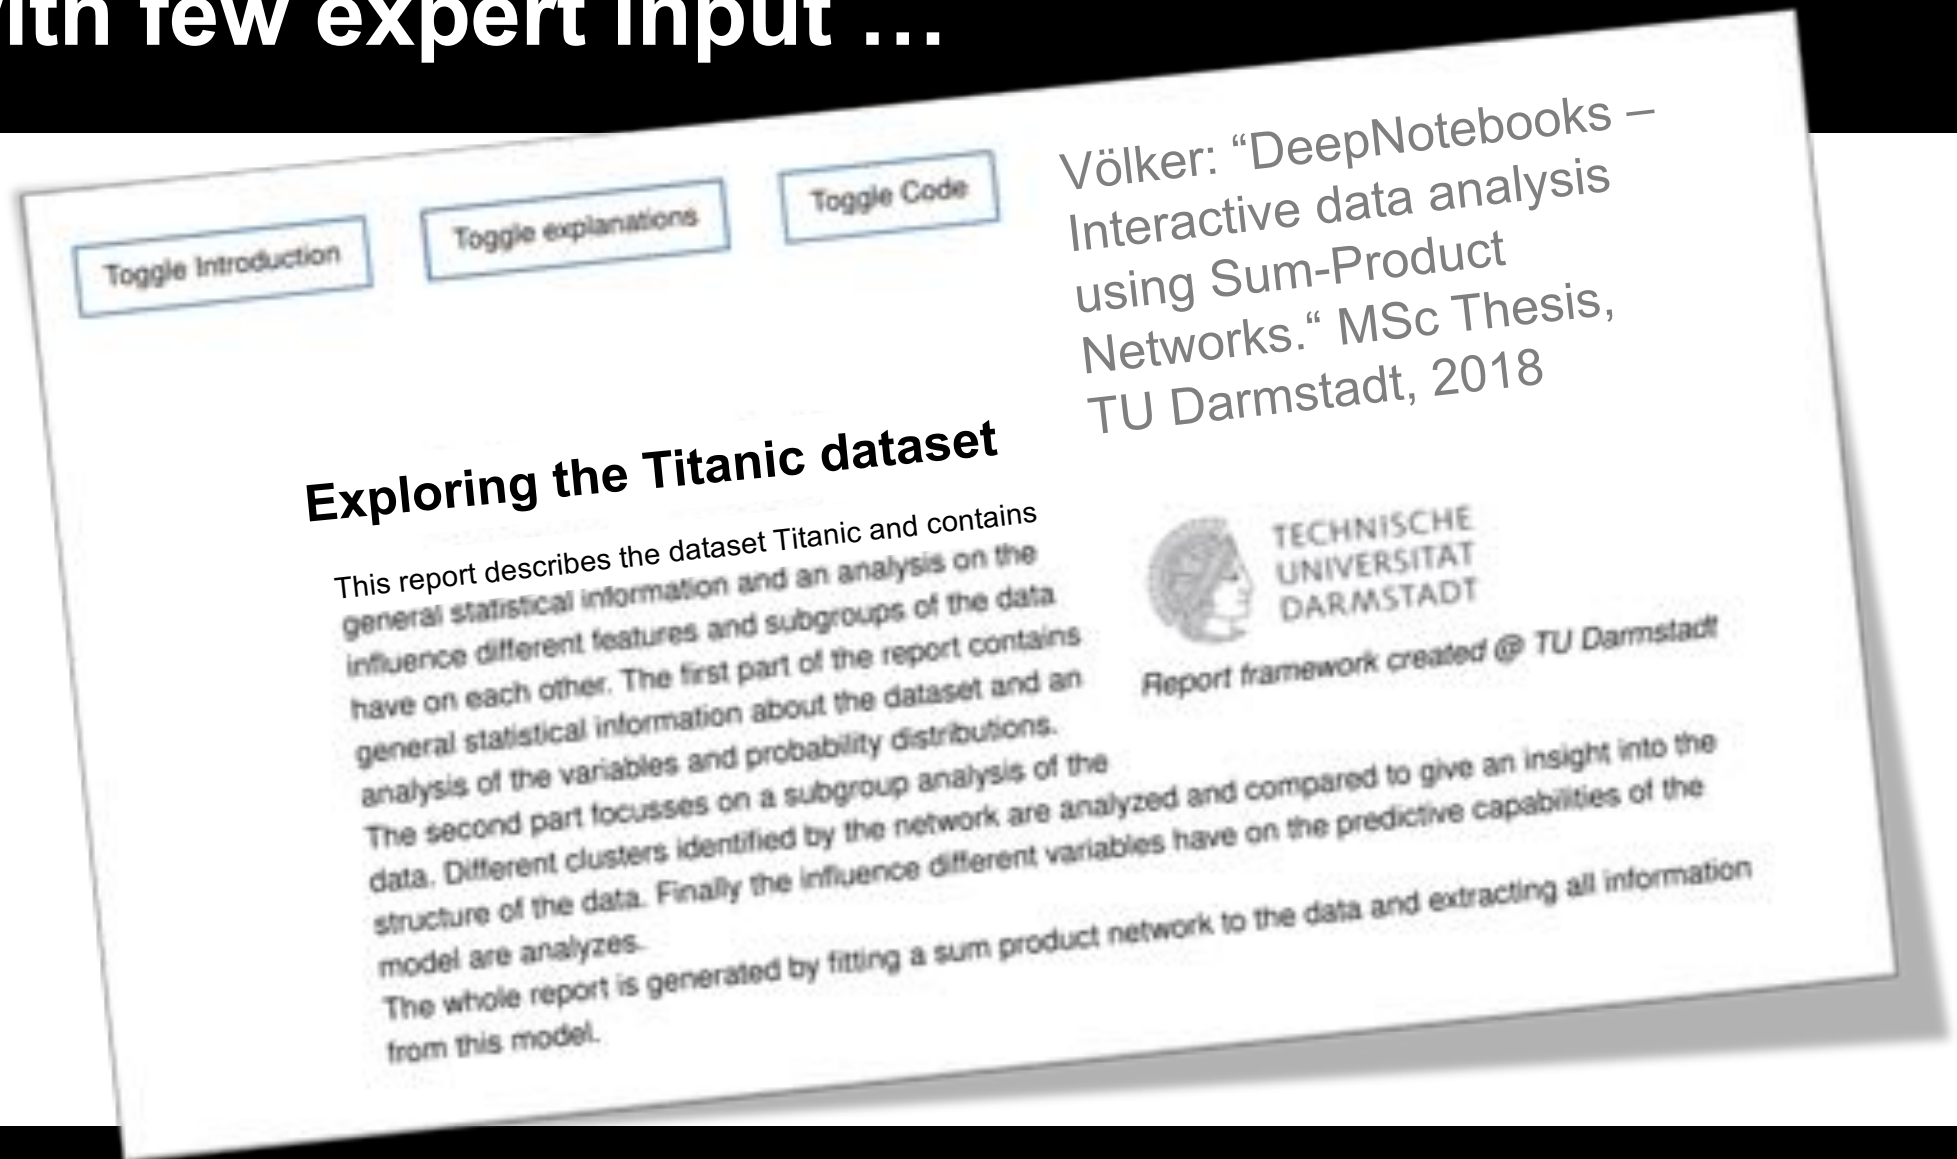
**Type-agnostic deep
probabilistic learning**

Automatic Bayesian Density Analysis



... can automatically discovers the statistical types and parametric forms of the variables

The machine understands the data with few expert input ...

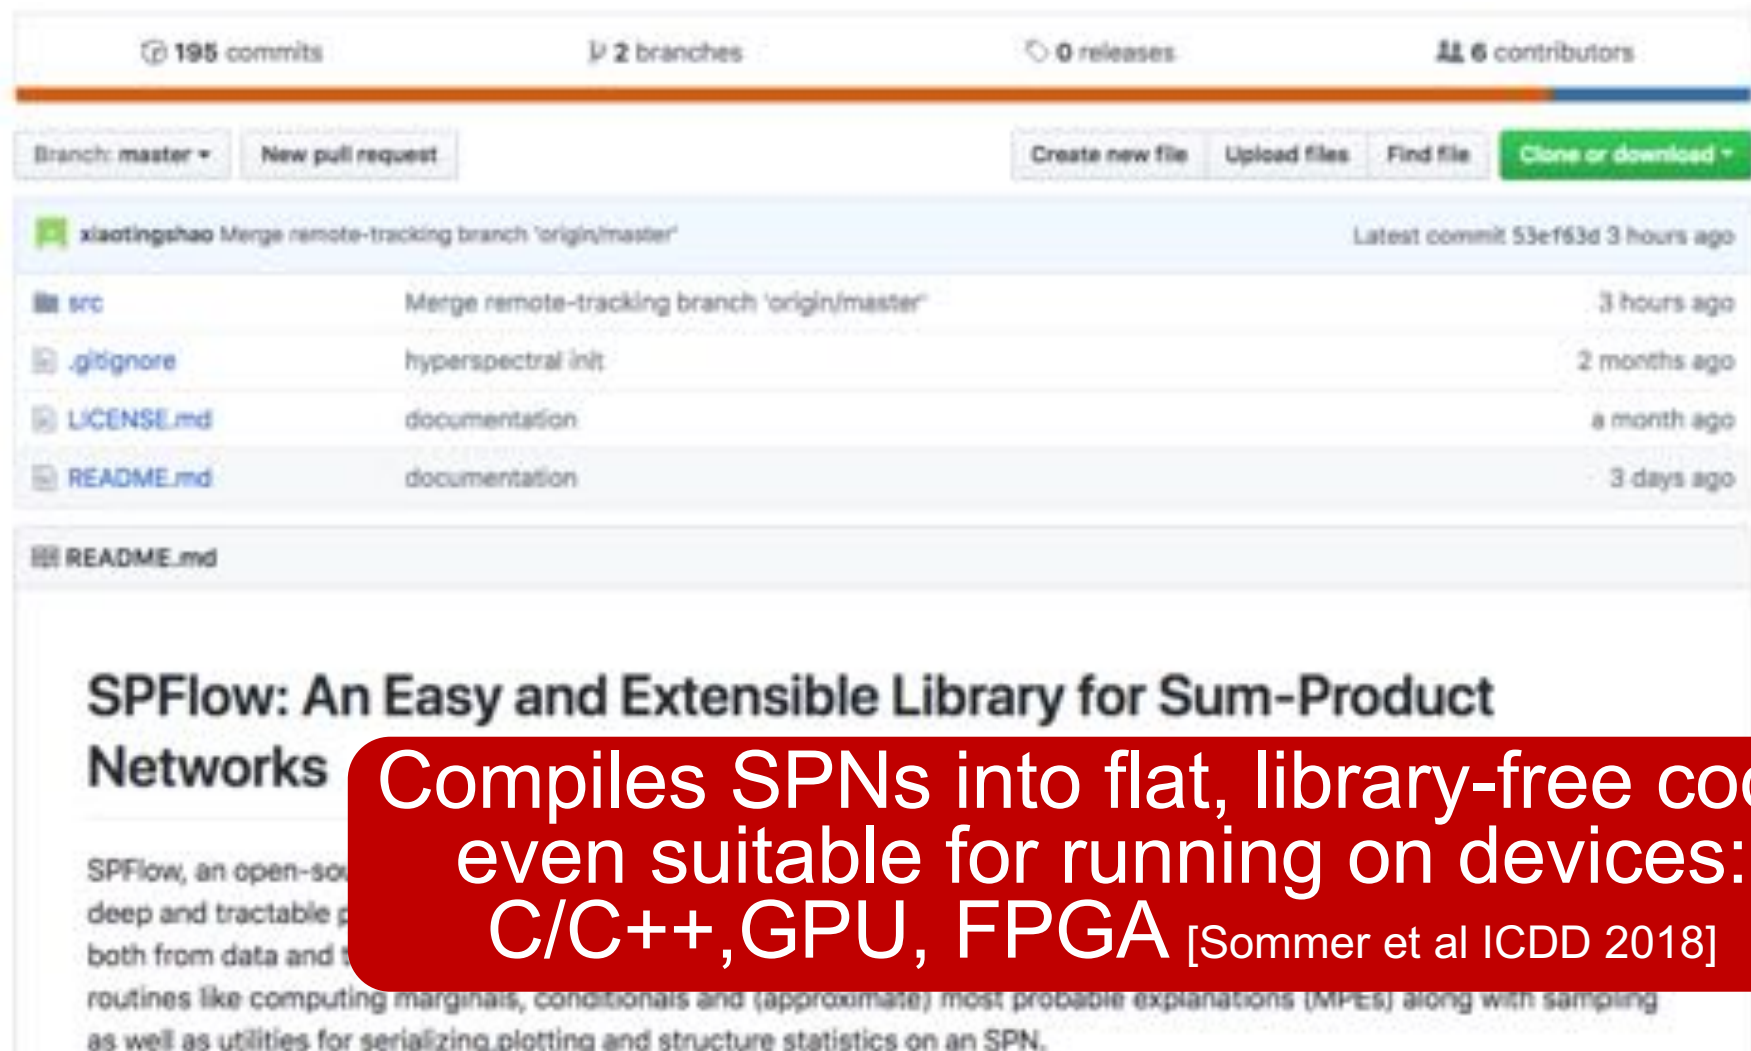


...and can compile data reports automatically

SPFlow: An Easy and Extensible Library for Sum-Product Networks

[Molina, Vergari, Stelzner, Peharz, Kersting 2018]

<https://github.com/alejandromolinaml/SPFlow>



The screenshot shows the GitHub repository for SPFlow. At the top, it displays repository statistics: 195 commits, 2 branches, 0 releases, and 6 contributors. Below this, there are buttons for 'Branch: master', 'New pull request', 'Create new file', 'Upload files', 'Find file', and a green 'Clone or download' button. A recent commit by xiaotingshao is shown, titled 'Merge remote-tracking branch 'origin/master'', with the latest commit hash 53ef63d and a timestamp of 3 hours ago. Below the commit list, there is a table of files in the repository:

File	Description	Last Commit
src	Merge remote-tracking branch 'origin/master'	3 hours ago
.gitignore	hyperspectral init	2 months ago
LICENSE.md	documentation	a month ago
README.md	documentation	3 days ago

Below the file list, the README.md content is partially visible, showing the title 'SPFlow: An Easy and Extensible Library for Sum-Product Networks' and a description of the library's capabilities.

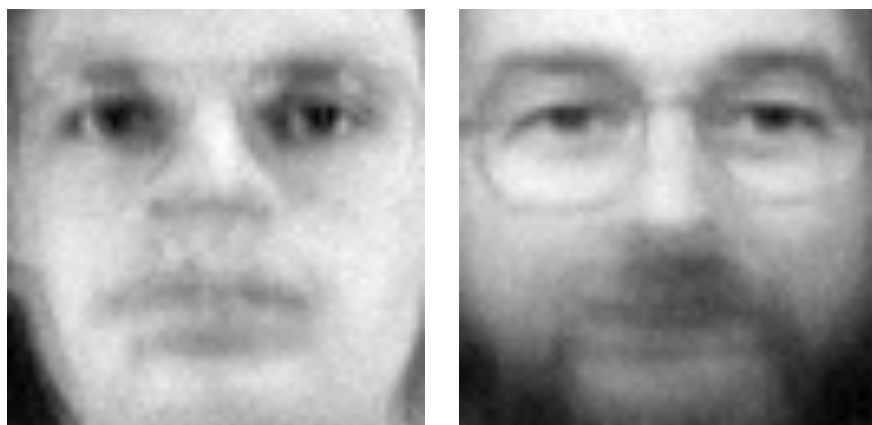
SPFlow: An Easy and Extensible Library for Sum-Product Networks

SPFlow, an open-source library, provides a deep and tractable model for Sum-Product Networks (SPNs) both from data and structure. It includes routines like computing marginals, conditionals and (approximate) most probable explanations (MPEs) along with sampling as well as utilities for serializing, plotting and structure statistics on an SPN.

Compiles SPNs into flat, library-free code even suitable for running on devices:
C/C++, GPU, FPGA [Sommer et al ICDD 2018]

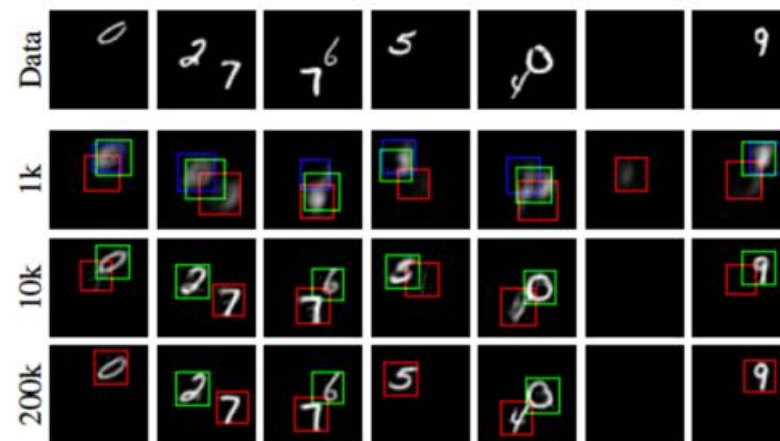
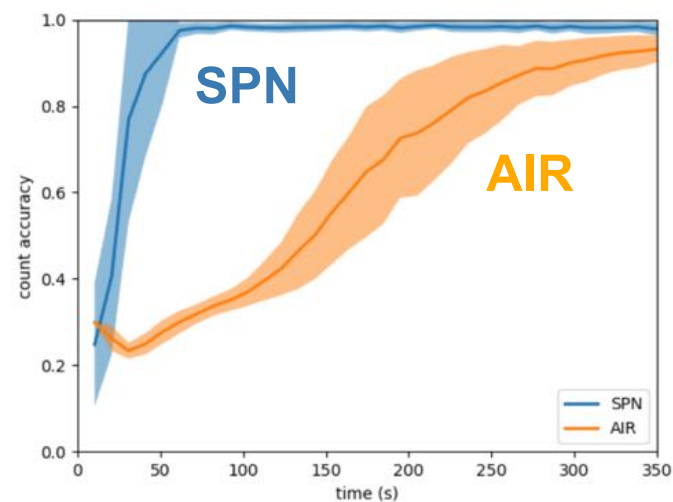
PixelSPNs

[Shao, Molina, Kersting 2018]



SPN AIR

[Stelzner, Peharz, Kersting 2018]



Next Steps



Symmetry-aware Deep Probabilistic Learning

Open-universe Mathematical Programming

Sum-Product Probabilistic Programming

Thanks

RelationalAI and
Apple, among others,
have invested hundreds
of millions of US dollars



And it appears in
industrial strength
solvers such as
CPLEX and GUROBI

