

Funsor: Functional Tensors for Probabilistic Programming

Fritz Obermeyer*, Eli Bingham*, Martin Jankowiak*, JP Chen, Du Phan



BROAD
INSTITUTE

Introduction: We need ‘autograd for integrals’

Probabilistic modelling and inference offer **a unifying approach to many machine learning tasks**, including quantifying uncertainty, learning structured generative models, producing interpretable explanations of data, and learning from weak or missing labels.

Probabilistic programming languages like Pyro allow **specification of probabilistic models in high-level programming languages**.

But many models that mix many different discrete and continuous variables still need custom inference strategies, and there is **no lower-level analogue of automatic differentiation software** intermediate between fully symbolic and fully black box integration.

Functional Tensors: a language for automatic integration over array-valued variables

Probabilistic programs generate lazy expressions with free variables. Inference algorithms integrate over free variables:

fun GenerativeModel(x)	$p \leftarrow 1$
$z \leftarrow \text{sample}(P_z)$	$p \leftarrow p \times P_z[v = z]$
$y \leftarrow \text{exp}(z)$	
observe($P_x[\theta = y], x$)	$p \leftarrow p \times P_x[\theta = y, v = x]$
end	maximize: $\sum_z p$

Approximate inference computations also generate lazy sum-product expressions in the same expression language.

fun GenerativeModel(x)	$p \leftarrow 1$
$z \leftarrow \text{sample}(P_z)$	$p \leftarrow p \times P_z[v = z]$
observe($P_x[\theta = z], x$)	$p \leftarrow p \times P_x[v = x, \theta = z]$
end	
fun InferenceModel(x)	$q \leftarrow 1$
$z \leftarrow \text{sample}(Q[\theta = x])$	$q \leftarrow q \times Q[v = z, \theta = x]$
end	maximize: $\sum_z q \log(p/q)$

Pyro represents terms with discrete free variables as torch.Tensors. We extend this representation to other functions, encoding them as “tensors” where some of the “dimensions” have size “real”:

$\tau \in \text{Type} ::= \mathbb{Z}_n$	“bounded integer”
$\mid \mathbb{Z}_{n_1} \times \cdots \times \mathbb{Z}_{n_k} \rightarrow \mathbb{R}$	“real-valued array”

We define a set of specific terms that are closed under variable substitution, sum-product operations, and various approximations:

$e \in \text{Funsor} ::= \text{Tensor}(\Gamma, w)$	“discrete factor”
$\mid \text{Gaussian}(\Gamma, i, P)$	“Gaussian factor”
$\mid \text{Delta}(v, e)$	“point mass”
$\mid \text{Variable}(v, \tau)$	“delayed value”
$\mid \widehat{f}(e_1, \dots, e_n)$	“apply function”
$\mid e_1[v = e_2]$	“substitute”
$\mid \sum_v e$	“marginalize”
$\mid \prod_{v/s} e$	“Markov product”

We extend lazy tensor expressions to include **dimensions of size ‘real’** and implement **semisymbolic integration**



On GitHub:
pyro-ppl/funsor

Approximation and transformation via (re-)interepretation

Most integrals cannot be computed directly and must be simplified. We rewrite lazy expressions by evaluating them with many different interpreters.

Some rules trigger PyTorch ops:

```
@eager.register(Binary, Op, Tensor, Tensor)
def eager_binary_tensor(op, lhs, rhs):
    inputs, (x, y) = align_tensors(lhs, rhs)
    data = op(x.data, y.data)
    return Tensor(data, inputs, lhs.dtype)
```

Some trigger further rewrites:

```
@eager.register(Binary, AddOp, Delta, Funsor)
def eager_add_delta(op, lhs, rhs):
    if lhs.name in rhs.inputs:
        rhs = rhs(**{lhs.name: lhs.point})
        return op(lhs, rhs)
    return None # defer to default implementation
```

Some rewrite subexpressions into approximate versions. monte_carlo rewrites Tensor and Gaussian to Delta:

```
@dispatched_interpretation
def monte_carlo(cls, *args): ...

@monte_carlo.register(Integrate, Funsor, Funsor, set)
def monte_carlo_integrate(log_measure, integrand, vs):
    log_measure = log_measure.sample()
    return eager.dispatch(
        Integrate, log_measure, integrand, vs)
```

Funsor expressions are closed under these approximation rewrites.

Example: detecting EEG changepoints

We use a moment matching interpreter to fit a switching linear dynamical system:

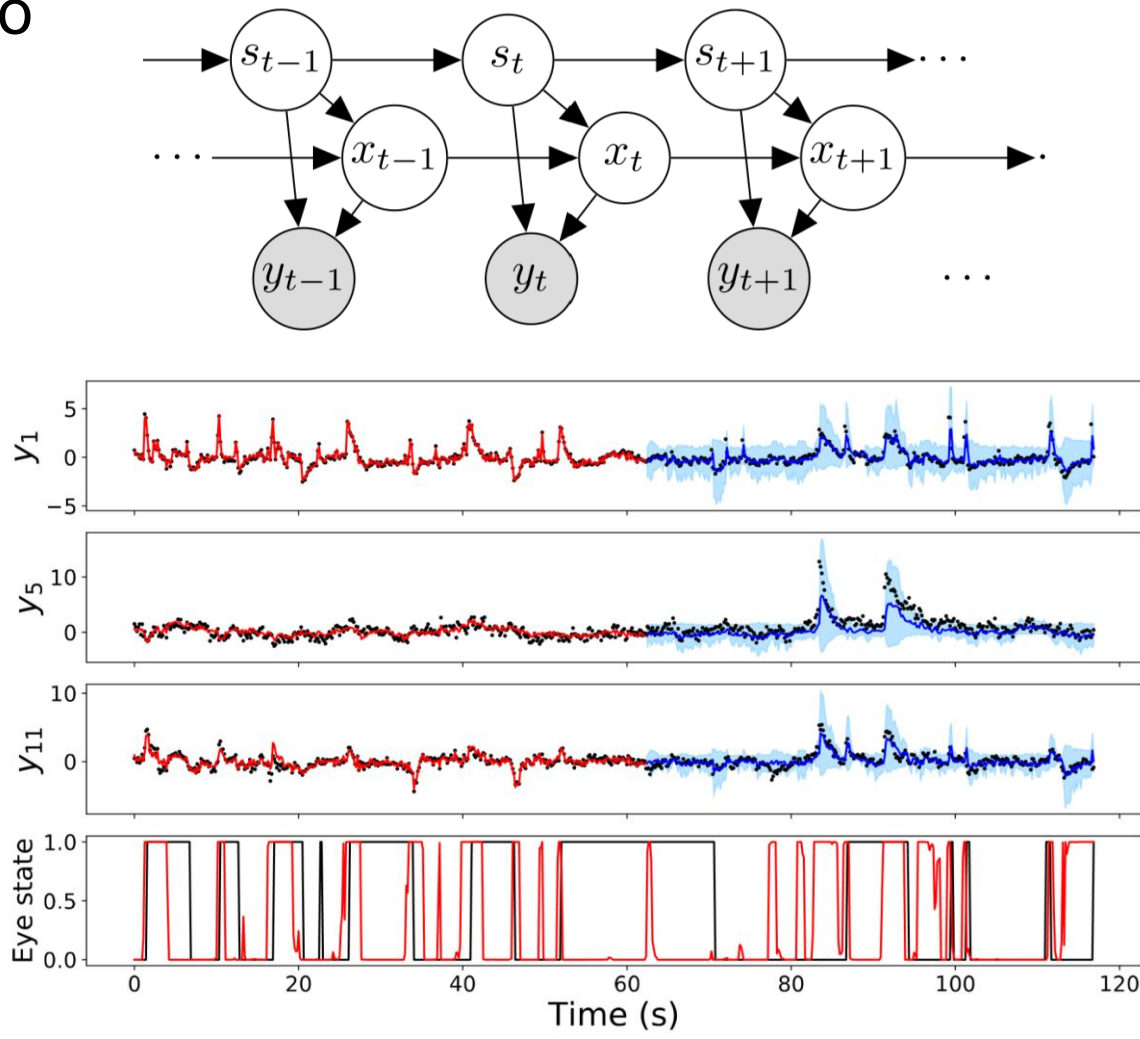
```
@interpretation(moment_matching_interpretation)
def marginal_log_prob(data, s_trans, x_trans, y_dist):
    ...
    for t, y in enumerate(data):
        ss[t] = Variable(f"s_{t}", bint(2))
        xs[t] = Variable(f"x_{t}", reals(5))
        ...
        log_prob += fdist.Categorical(
            s_trans(s=ss[t - 1]), value=ss[t])
        log_prob += x_trans(s=ss[t], x=xs[t - 1], y=xs[t])

    log_prob = log_prob.reduce(ops.logaddexp,
        {ss[t - 2].name, xs[t - 2].name})

    log_prob += y_dist(s=s_vars[t], x=x_vars[t], y=y)

    for t in range(2):
        log_prob.reduce(ops.logaddexp,
            {s_vars[T - 2 + t].name, x_vars[T - 2 + t].name})

    return log_prob
```



Extending the language: parallel-scan over sequential structure

Many common sum-product expressions have linear chain structures. We define a generic operation on atomic funsor terms for collapsing this structure:


For Tensor terms, the MarkovProduct op corresponds to chain matrix multiplication:

$$\prod_{t/(i,j)} f = f[t=0] \bullet f[t=1] \bullet \cdots \bullet f[t=T-1]$$


i.e. the binary operation is a GEMM:

$$f \bullet g = \sum_k f[j=k] \times g[i=k]$$

Our parallel-scan algorithm computes this in **O(log(T))** on a T-processor parallel machine:






Algorithm 1 MARKOVPRODUCT
input a funsor f , a time variable $t \in \text{fv}(f)$, a step mapping $s \subseteq \text{fv}(f) \times \text{fv}(f)$.
output the Markov product funsor $\prod_{t/s} f$.
Create substitutions with fresh names (barred):
 $s_e \leftarrow \{(y, \bar{x}) \mid (x, y) \in s\}$ to rename even factors, and
 $s_o \leftarrow \{(x, \bar{x}) \mid (x, y) \in s\}$ to rename odd factors.
Let $v \leftarrow \{\bar{x} \mid (x, y) \in s\}$ be variables to marginalize.
Let $T \leftarrow |\Gamma_f[t]|$ be the length of the time axis.
while $T > 1$ **do**
 Split f into even and odd parts of equal length:
 $f_e \leftarrow f[s_e, t = (0, 2, 4, 6, \dots, 2\lfloor T/2 \rfloor - 2)]$
 $f_o \leftarrow f[s_o, t = (1, 3, 5, 7, \dots, 2\lfloor T/2 \rfloor - 1)]$
 Perform parallel sum-product contraction:
 $f' \leftarrow \sum_v f_e \times f_o$
 if T is even **then** $f \leftarrow f'$
 else $f \leftarrow \text{concat}_t(f', f[t = T - 1])$;
 $T \leftarrow \lceil T/2 \rceil$
return $f[t = 0]$



Example: forecasting BART ridership

We do collapsed variational inference in a neural Kalman filter to model rides between all 47 BART stations for 10 years of hour-level counts:



```
def model(features, counts):
    gr_t = funsor.Variable(
        "gr_t", ...)[ "time" ]
    init, trans, obs, rate = ...
    ...
    prior = MarkovProduct(
        ops.logaddexp, ops.add,
        trans + obs(gate_rate=gr_t),
        "time", {"state": "state(time=1)"})

    llk = fdist.Poisson(rate["origin", "destin"])
    return prior + llk(value=counts)
```

```
def guide(features, counts):
    loc, sc = ...
    diag_normal = fdist.Normal(
        loc, sc, value="gr_t")
    return diag_normal
```

```
def elbo_loss(features, counts):
    q = guide(features, counts)

    with interpretation(lazy):
        p = model(features, counts)
        pq = p - q

    with interpretation(monte_carlo):
        elbo = funsor.Integrate(q, pq)

    return elbo
```