

TL;DR

- Rejection sampling loops are widely used in implementing probabilistic models e.g. sampling from constrained distributions.
- Unbounded loops (including rejection sampling) add major complexities to the inference task.
- We address the problem of efficient amortized importance-sampling-based inference, in particular, Inference Compilation (IC) [4] in such models.
- We show naive application of IC can produce estimators with unbounded variance.
- We propose an unbiased estimator for handling rejection sampling loop, prove its finite variance and implement it in pyprob [1; 2] in a way that requires minimal modifications to user code.

<pre> 1: $x \sim p(x)$ 2: 3: for $k \in \mathbb{N}^+$ do 4: $\mathbf{z}^k \sim p(\mathbf{z} x)$ 5: 6: 7: if $c(x, \mathbf{z}^k)$ then 8: $\mathbf{z} = \mathbf{z}^k$ 9: break 10: observe($y, p(y \mathbf{z}, x)$) (a) Original program </pre>	<pre> $x \sim q(x y)$ $w \leftarrow \frac{p(x)}{q(x y)}$ for $k \in \mathbb{N}^+$ do $\mathbf{z}^k \sim q(\mathbf{z} x, y)$ $w^k \leftarrow \frac{p(\mathbf{z}^k x)}{q(\mathbf{z}^k x, y)}$ $w \leftarrow w w^k$ if $c(x, \mathbf{z}^k)$ then $\mathbf{z} = \mathbf{z}^k$ break $w \leftarrow w p(y \mathbf{z}, x)$ (b) Inference compilation </pre>
<pre> 1: $x \sim p(x)$ 2: 3: $\mathbf{z} \sim p(\mathbf{z} x, c(x, \mathbf{z}))$ 4: 5: observe($y, p(y \mathbf{z}, x)$) (c) Equivalent to above </pre>	<pre> $x \sim q(x y)$ $w \leftarrow \frac{p(x)}{q(x y)}$ $\mathbf{z} \sim q(\mathbf{z} x, y, c(x, \mathbf{z}))$ $w \leftarrow w \frac{p(\mathbf{z} x, c(x, \mathbf{z}))}{q(\mathbf{z} x, y, c(x, \mathbf{z}))}$ $w \leftarrow w p(y \mathbf{z}, x)$ (d) Our IS estimator </pre>

IC weighting

$$w_{IC} = \frac{p(x)}{q(x|y)} p(y|x, z) \prod_{k=1}^L w^k$$

Theorem: Under some mild conditions if the following holds then the variance of w_{IC} is infinite.

$$\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z}|x, y)} \left[\frac{p(\mathbf{z}|x)^2}{q(\mathbf{z}|x, y)^2} (1 - p(A|x, \mathbf{z})) \right] \geq 1$$

where A is the event of $c(x, \mathbf{z})$ being satisfied.

Collapsed weighting

$$w_C = \frac{p(x)}{q(x|y)} \frac{p(\mathbf{z}|x, A)}{q(\mathbf{z}|x, y, A)} p(y|x, \mathbf{z})$$

- $\mathbb{E}[w_{IC}] = \mathbb{E}[w_C]$ but this weighting scheme does not introduce infinite variance to the estimator.
- Unfortunately, we cannot directly compute w_C

Our method (ARS)

$$w_C = \frac{p(x)}{q(x|y)} \frac{p(\mathbf{z}|x)}{q(\mathbf{z}|x, y)} p(y|x, \mathbf{z}) \frac{q(A|x, y)}{p(A|x)}$$

- $q(A|x, y)$ is a Bernoulli distribution. We get an unbiased estimate of its parameter via Monte Carlo.
- The number of samples $p(\mathbf{z}|x)$ until the first acceptance follows a geometric distribution with mean $\frac{1}{p(A|x)}$.
- We get an unbiased Monte Carlo estimates of $q(A|x, y)$ and $\frac{1}{p(A|x)}$.

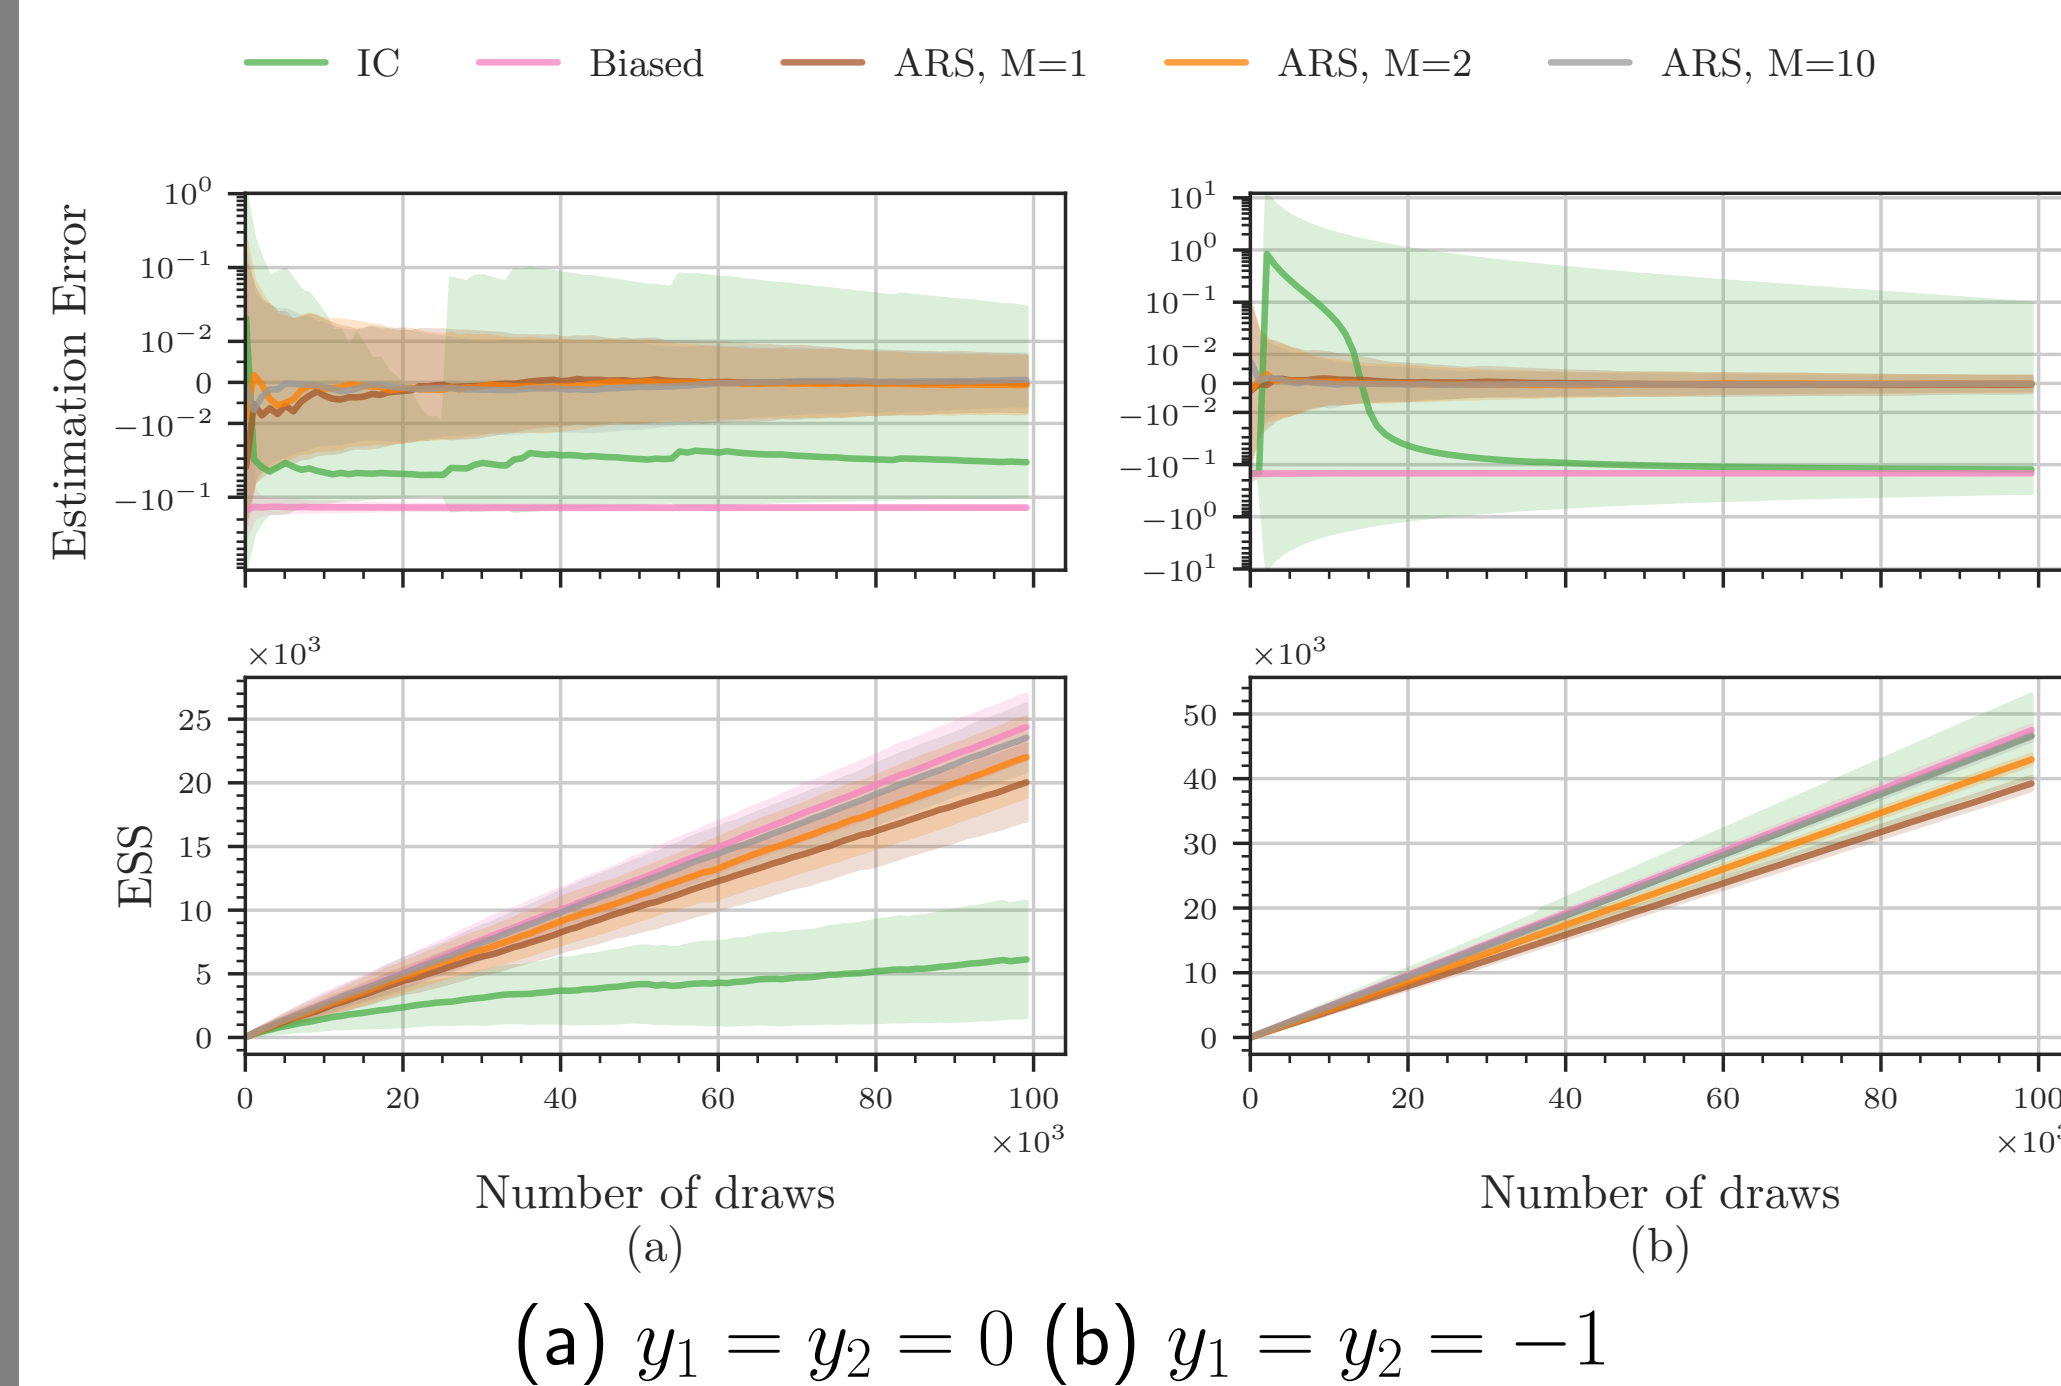
Marsaglia Experiment

Marsaglia polar method [5] is a pseudo-random number sampling method for generating samples from a Normal distribution.

$$(a, b) \sim \text{Unit circle} \quad \mu = a \sqrt{\frac{-2 \log(a^2 + b^2)}{a^2 + b^2}} \equiv \mu \sim \mathcal{N}(0, 1)$$

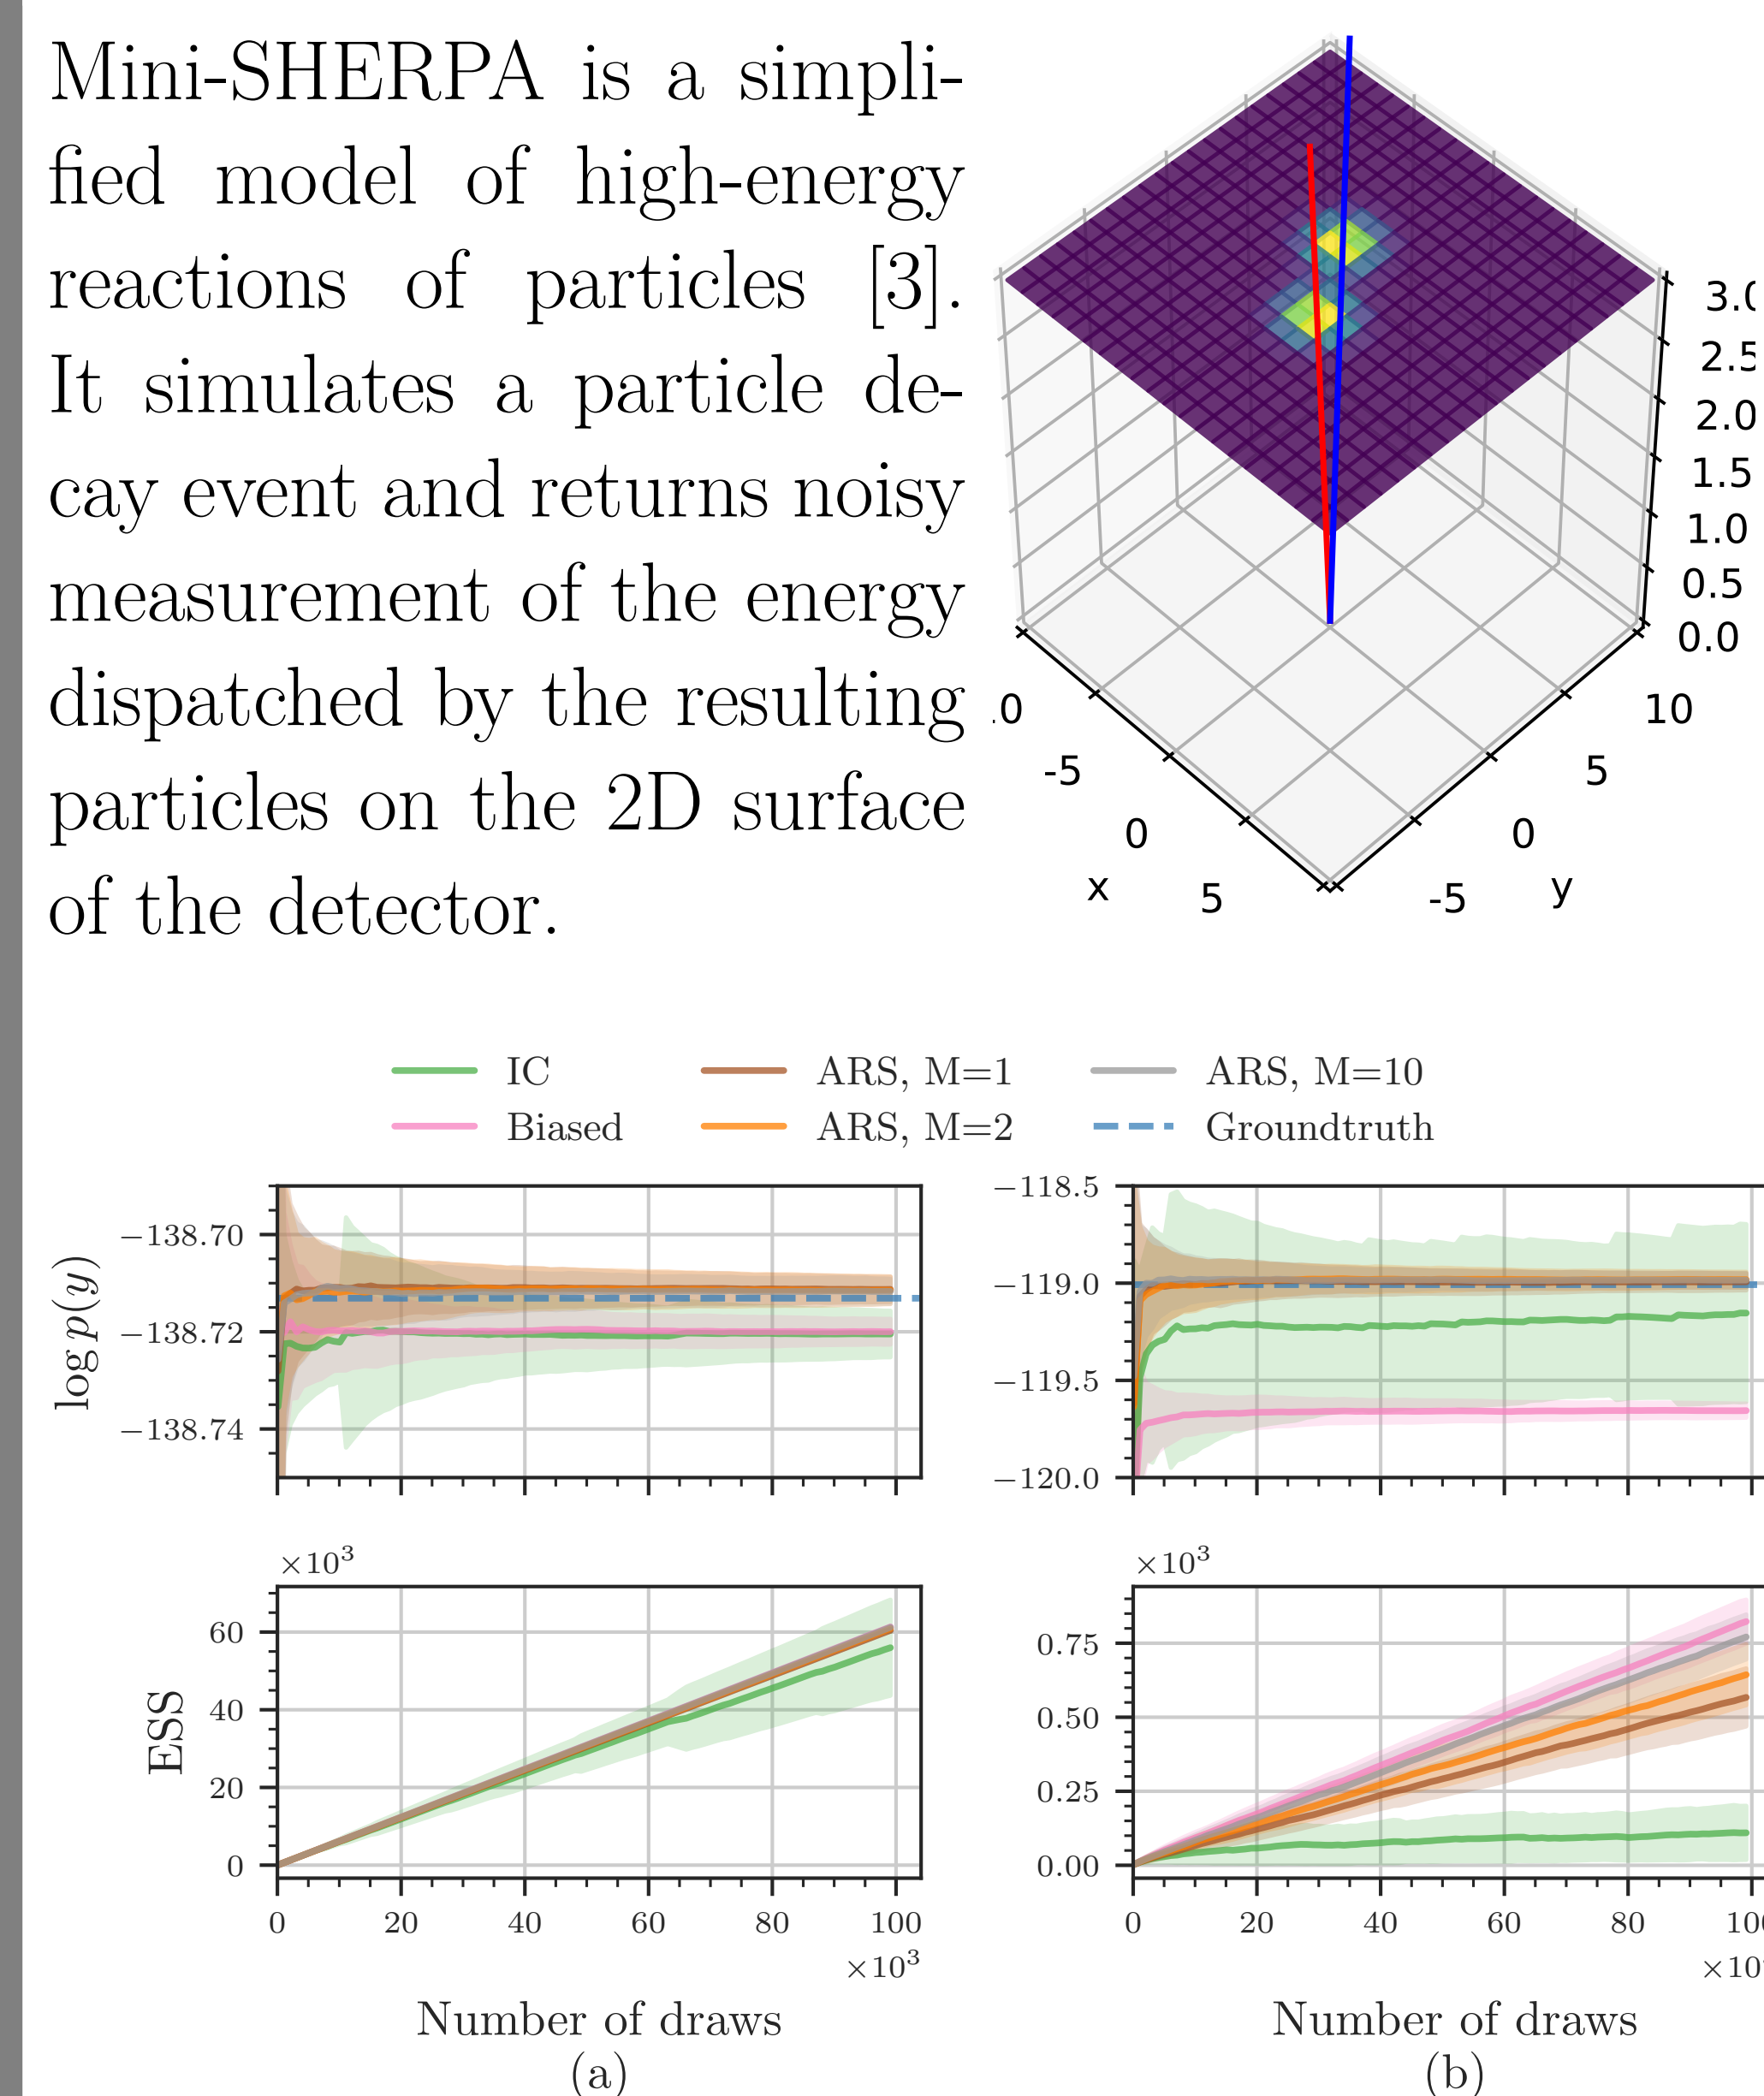
$$y_i \sim \mathcal{N}(\mu, \sigma^2), i \in \{1, 2\}$$

We target estimating $p(y_1, y_2)$.



Mini-SHERPA Experiment

Mini-SHERPA is a simplified model of high-energy reactions of particles [3]. It simulates a particle decay event and returns noisy measurement of the energy dispatched by the resulting particles on the 2D surface of the detector.



Algorithm

```

1:  $x \sim q(x|y)$ 
2:  $w \leftarrow \frac{p(x)}{q(x|y)}$ 
3: for  $k \in \mathbb{N}^+$  do
4:    $\mathbf{z}^k \sim q(\mathbf{z}|x, y)$ 
5:   if  $c(x, \mathbf{z}^k)$  then
6:      $\mathbf{z} = \mathbf{z}^k$ 
7:   break
8:  $w \leftarrow w \frac{p(\mathbf{z}|x)}{q(\mathbf{z}|x, y)}$ 
9:  $K \leftarrow 0$ 
10: for  $i \in 1, \dots, N$  do
11:    $\mathbf{z}'_i \leftarrow q(\mathbf{z}|x, y)$ 
12:    $K \leftarrow K + c(\mathbf{z}, x)$ 
13: for  $j \in 1, \dots, M$  do
14:   for  $l \in \mathbb{N}^+$  do
15:      $\mathbf{z}''_{j,l} \leftarrow q(\mathbf{z}|x, y)$ 
16:     if  $c(x, \mathbf{z}''_{j,l})$  then
17:        $T_j \leftarrow l$ 
18:     break
19:  $T \leftarrow \frac{1}{M} \sum_{j=1}^M T_j$ 
20:  $w \leftarrow w \frac{KT}{N}$ 
21:  $w \leftarrow w p(y|\mathbf{z}, x)$ 

```

Implementation

We introduce two new functions to tag the beginning and end of rejection sampling loops.

Original	Annotated
<pre> x = sample(P.x) while True: z = sample(P.z(x)) if c(x, z): break observe(P.y(x, z), y) return x, z </pre>	<pre> x = sample(P.x) while True: rs.start() z = sample(P.z(x)) if c(x, z): rs.end() break observe(P.y(x, z), y) return x, z </pre>

References

- A. G. Baydin, L. Heinrich, W. Bhimji, B. Gram-Hansen, G. Louppe, L. Shao, K. Cranmer, F. Wood, et al. Efficient probabilistic inference in the quest for physics beyond the standard model. In *Thirty-second Conference on Neural Information Processing Systems (NeurIPS)*, 2019.
- A. G. Baydin, L. Shao, W. Bhimji, L. Heinrich, L. Meadows, J. Liu, A. Munk, S. Naderiparizi, B. Gram-Hansen, G. Louppe, et al. Etalumis: Bringing probabilistic programming to scientific simulators at scale. In *the International Conference for High Performance Computing, Networking, Storage and Analysis (SC '19)*, 2019.
- T. Gleisberg, S. Höche, F. Krauss, M. Schönherr, S. Schumann, F. Siegert, and J. Winter. Event generation with sherpa 1.1. *Journal of High Energy Physics*, 2009(02):007, 2009.
- T. A. Le, A. G. Baydin, and F. Wood. Inference compilation and universal probabilistic programming. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, volume 54 of *Proceedings of Machine Learning Research*, pages 1338–1348, Fort Lauderdale, FL, USA, 2017. PMLR.
- G. Marsaglia and T. A. Bray. A convenient method for generating normal variables. *SIAM review*, 6(3):260–264, 1964.